

Ovarian Cancer Histotypes: Report of Statistical Findings

Derek Chiu

May 22, 2025

Table of contents

Preface	6
1 Introduction	7
2 Methods	8
2.1 Pre-Processing	8
2.1.1 Case Selection	8
2.1.2 Quality Control	8
2.1.3 Housekeeping Genes Normalization	9
2.1.4 Between CodeSet and Site Normalization	9
2.1.5 Final Processing	10
2.2 Classifiers	11
2.2.1 Resampling of Training Set	12
2.2.2 Hyperparameter Tuning	12
2.2.3 Subsampling	12
2.2.4 Workflows	13
2.3 Two-Step Algorithm	13
2.3.1 Aggregating Predictions	14
2.4 Sequential Algorithm	15
2.4.1 Aggregating Predictions	16
2.5 Performance Evaluation	17
2.5.1 Class Metrics	17
2.5.2 AUC	19
2.6 Rank Aggregation	19
2.7 Gene Optimization	19
2.7.1 Variable Importance	21
3 Distributions	23
3.1 Histotype Distribution	23
3.2 Cohort Distribution	25
3.3 Quality Control	25
3.3.1 Failed Samples	25
3.3.2 %GD vs. SNR	27
3.4 Pairwise Gene Expression	29
4 Results	33
4.1 Training Set	34
4.1.1 Accuracy	34
4.1.2 Sensitivity	36
4.1.3 Specificity	38
4.1.4 F1-Score	40

4.1.5	Balanced Accuracy	42
4.1.6	Kappa	44
4.1.7	G-mean	45
4.2	Rank Aggregation	47
4.2.1	Across Classes	48
4.2.2	Across Metrics	51
4.2.3	Top Workflows	51
4.3	Optimal Gene Sets	54
4.3.1	Sequential Algorithm	54
4.3.2	SMOTE-Random Forest	57
4.3.3	Two-Step	60
4.4	Test Set Performance	62
4.4.1	Confirmation Set	65
4.4.2	Validation Set	73

References	77
-------------------	-----------

List of Figures

2.1	Venn diagram of common and unique gene targets covered by each CodeSet	10
2.2	Cohorts Selection	11
2.3	Visualization of Subsampling Techniques	13
2.4	Two-Step Algorithm	14
2.5	Aggregating Predictions for Two-Step Algorithm	15
2.6	Sequential Algorithm	16
2.7	Aggregating Predictions for Sequential Algorithm	17
3.1	% Genes Detected vs. Signal to Noise Ratio	27
3.2	% Genes Detected vs. Signal to Noise Ratio (Zoomed)	28
3.3	Random1-Normalized CS1 vs. CS3 Gene Expression	29
3.4	Random1-Normalized CS2 vs. CS3 Gene Expression	30
3.5	HKgenes-Normalized CS1 vs. CS3 Gene Expression	31
3.6	HKgenes-Normalized CS2 vs. CS3 Gene Expression	32
4.1	Training Set Mean Accuracy	35
4.2	Training Set Mean Sensitivity	37
4.3	Training Set Mean Specificity	39
4.4	Training Set Mean F1-Score	41
4.5	Training Set Mean Balanced Accuracy	43
4.6	Training Set Mean Kappa	45
4.7	Training Set Mean G-mean	47
4.8	Top 5 Workflow Per-Class Evaluation Metrics by Metric	52
4.9	Top 5 Workflow Per-Class Evaluation Metrics by Metric	53
4.10	Gene Optimization for Sequential Classifier	54
4.11	Gene Optimization for SMOTE-Random Forest Classifier	57
4.12	Gene Optimization for Two-Step Classifier	60
4.13	Gene Optimized Workflows Per-Class Metrics in Confirmation Set	66
4.14	Confusion Matrices for Confirmation Set Models	67
4.15	ROC Curves for Sequential Full Model in Confirmation Set	68
4.16	ROC Curves for Sequential, Optimal Model in Confirmation Set	69
4.17	ROC Curves for SMOTE-Random Forest, Full Set Model in Confirmation Set	70
4.18	ROC Curves for SMOTE-Random Forest, Optimal Set Model in Confirmation Set	71
4.19	ROC Curves for Two-Step Full Model in Confirmation Set	72
4.20	ROC Curves for Two-Step Optimal Model in Confirmation Set	73
4.21	SMOTE-Random Forest Per-Class Metrics in Validation Set	74
4.22	Confusion Matrix for Validation Set Model	74
4.23	ROC Curves for SMOTE-Random Forest, Optimal Set Model in Validation Set	75
4.24	Subtype Prediction Summary among Predicted HGSC Samples	76

List of Tables

2.1	Gene Distribution	20
3.1	Histotype Distribution in Training Set by Processing Stage	23
3.2	Histotype Distribution in Training, Confirmation, and Validation Sets	24
3.3	Pre-QC Cohort Distribution by CodeSet	25
3.4	Quality Control Summary	26
4.1	Training Set Mean Accuracy	34
4.2	Training Set Mean Sensitivity	36
4.3	Training Set Mean Specificity	38
4.4	Training Set Mean F1-Score	40
4.5	Training Set Mean Balanced Accuracy	42
4.6	Training Set Mean Kappa	44
4.7	Training Set Mean G-mean	46
4.8	F1-Score Rank Aggregation Summary	48
4.9	Balanced Accuracy Rank Aggregation Summary	49
4.10	Kappa Rank Aggregation Summary	50
4.11	Rank Aggregation Comparison of Metrics Used	51
4.12	Top 5 Workflows from Final Rank Aggregation	51
4.13	Gene Profile of Optimal Set in Sequential Algorithm	54
4.14	Gene Profile of Optimal Set in SMOTE-Random Forest Workflow	57
4.15	Gene Profile of Optimal Set in Two-Step Workflow	60
4.16	Evaluation Metrics on Confirmation Set Models	65
4.17	Evaluation Metrics on Validation Set Model, SMOTE-Random Forest, Optimal Set	73

Preface

This report of statistical findings describes the classification of ovarian cancer histotypes using data from NanoString CodeSets.

Marina Pavanello conducted the initial exploratory data analysis, Cathy Tang implemented class imbalance techniques, Derek Chiu conducted the normalization and statistical analysis, and Lauren Tindale and Aline Talhouk are the project leads.

1 Introduction

Ovarian cancer has five major histotypes: high-grade serous carcinoma (HGSC), low-grade serous carcinoma (LGSC), endometrioid carcinoma (ENOC), mucinous carcinoma (MUC), and clear cell carcinoma (CCOC). A common problem with classifying these histotypes is that there is a class imbalance issue. HGSC dominates the distribution, commonly accounting for 70% of cases in many patient cohorts, while the other four histotypes are spread over the rest of the cases. Subsampling methods like up-sampling, down-sampling, and SMOTE can be used to mitigate this problem.

The supervised learning is performed under a consensus framework: we consider various classification algorithms and use evaluation metrics like accuracy, F1-score, and Kappa, to inform the decision of which methods to carry forward for prediction in confirmation and validation sets.

2 Methods

2.1 Pre-Processing

2.1.1 Case Selection

Prior to pre-processing, samples were split into a training, a confirmation, and a validation set.

- Training
 - CS1: MAYO, OOU, OOUE, VOA, MTL
 - CS2: MAYO, OOU, OOUE, OVAR3, VOA, ICON7, JAPAN, MTL, POOL-CTRL
 - CS3: OOU, OOUE, VOA, POOL-1, POOL-2, POOL-3
- Confirmation:
 - CS3: TNCO
- Validation:
 - CS3: DOVE4

2.1.2 Quality Control

Before normalization, we calculated several quality control measures and excluded samples that failed to achieve sample quality in one or more of these measures.

- **Linearity of positive control genes:** If the R-squared from a linear model of positive controls and their concentrations is less than 0.95 or missing, then the sample is flagged.
- **Imaging quality:** The sample is flagged if the field of view percentage is less than 75%.
- **Positive Control flag:** We consider the two smallest positive controls at concentrations 0.5 and 1. If these two controls are less than the lower limit of detection (defined as two standard deviations below the mean of the negative control expression), or if the mean negative control expression is 0, the sample is flagged.
- **The signal-to-noise ratio or percent of genes detected:** These two measures are defined as the ratio of the average housekeeping gene expression over the upper limit of detection, defined as two standard deviations above the mean of the negative control expression (or 0 if this limit is less than 0.001), and the proportion of endogenous genes with expression greater than the upper limit of detection. These measures are flagged if they are below a pre-specified threshold, which is determined visually by considering their bivariate distribution in a scatterplot. In this case, we used 100 for the SNR threshold and 50% for the threshold for genes detected. Note: these thresholds were determined by examining the relationship in [Section 3.3.2](#).

2.1.3 Housekeeping Genes Normalization

The full training set (n=1243) comprised of data from three CodeSets (CS) 1, 2, and 3. Data normalization removes technical variation from high-throughput platforms to improve the validity of comparative analyses.

Each CodeSet was first normalized to housekeeping genes: *ACTB*, *RPL19*, *POLR1B*, *SDHA*, and *PGK1*. Housekeeping genes encode proteins responsible for basic cell function and have consistent expression in all cells. All expression values were log2 transformed. Normalization to housekeeping genes corrects the viable RNA from each sample. This is achieved by subtracting the average log (base 2)-transformed expression of the housekeeping genes from the log (base 2)-transformed expression of each gene:

$$\log_2(\text{endogenous gene expression}) - \text{average}(\log_2(\text{housekeeping gene expression})) = \text{relative expression} \quad (2.1)$$

2.1.4 Between CodeSet and Site Normalization

To normalize between CodeSets, we randomly selected five specimens, one from each histotype, among specimens repeated in all three CodeSets. This formed the reference set (Random 1). We selected only one sample from each histotype to use as few samples as possible for normalization and retain the rest for analysis.

A reference-based approach (Talhouk et al. (2016)) was used to normalize CS1 to CS3 and CS2 to CS3 across their common genes:

$$\text{X-Norm}_{\text{CS1}} = X_{\text{CS1}} + \bar{R}_{\text{CS3}} - \bar{R}_{\text{CS1}} \quad \text{X-Norm}_{\text{CS2}} = X_{\text{CS2}} + \bar{R}_{\text{CS3}} - \bar{R}_{\text{CS2}} \quad (2.2)$$

Samples in CS3 were processed at three different locations; we also had to normalize for “site” in this CodeSet. Finally, the CS3 expression samples were included in the training set without further normalization:

$$\text{X-Norm}_{\text{CS3-USC}} = X_{\text{CS3-USC}} + \bar{R}_{\text{CS3-VAN}} - \bar{R}_{\text{CS3-USC}} \quad \text{X-Norm}_{\text{CS3-AOC}} = X_{\text{CS3-AOC}} + \bar{R}_{\text{CS3-VAN}} - \bar{R}_{\text{CS3-AOC}} \quad (2.3)$$

Finally, the CS3 expression samples were included in the training set without further normalization. The initial training set is assembled by combining all four of the previously mentioned normalized datasets along with the two CS3 expression subsets not used in normalization:

$$\begin{aligned} \text{Training Set} &= \text{X-Norm}_{\text{CS1}} + \text{X-Norm}_{\text{CS2}} + \text{X-Norm}_{\text{CS3-USC}} + \text{X-Norm}_{\text{CS3-AOC}} + \text{X-Norm}_{\text{CS3}} + \text{X-Norm}_{\text{CS3-VAN}} \\ &= \text{X-Norm}_{\text{CS1}} + \text{X-Norm}_{\text{CS2}} + \text{X-Norm}_{\text{CS3}} \end{aligned} \quad (2.4)$$



Figure 2.1: Venn diagram of common and unique gene targets covered by each CodeSet

2.1.5 Final Processing

We map ovarian histotypes to all remaining samples and keep the major histotypes for building the predictive model: high-grade serous carcinoma (HGSC), clear cell ovarian carcinoma (CCOC), endometrioid ovarian carcinoma (ENOC), low-grade serous carcinoma (LGSC), mucinous carcinoma (MUC).

Duplicate cases (two samples with the same ottaID) were removed before generating the final training set to use for fitting the classification models. All CS3 cases were preferred over CS1

and CS2, and CS3-Vancouver cases were preferred over CS3-AOC and CS3-USC when selecting duplicates.

The final training set used only genes that were common across all three CodeSets.

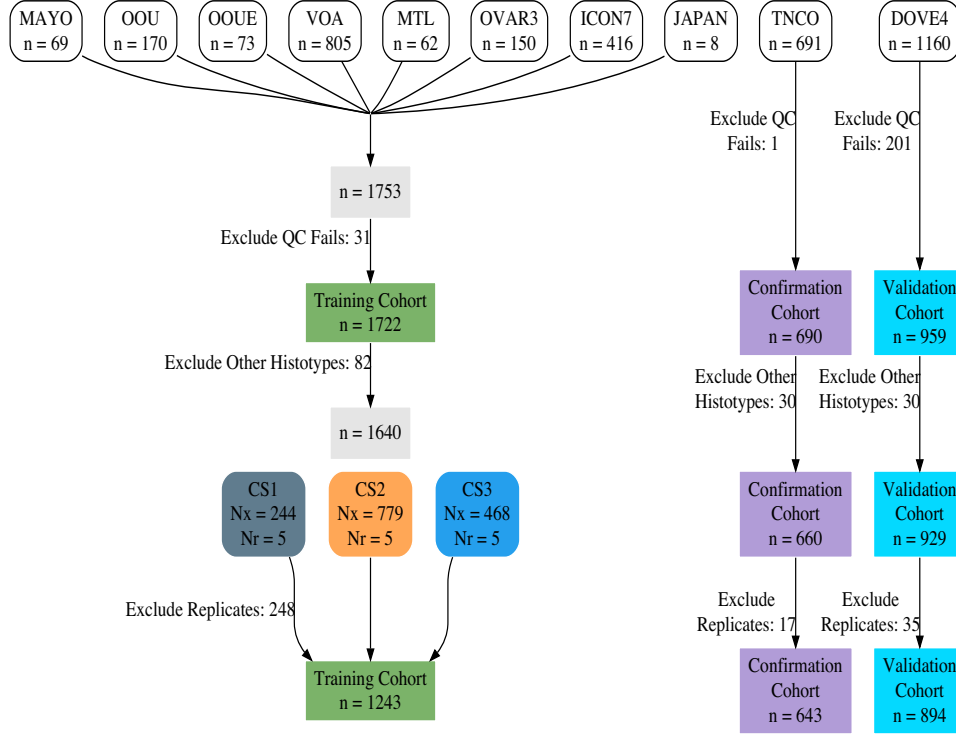


Figure 2.2: Cohorts Selection

2.2 Classifiers

We use 4 classification algorithms in the supervised learning framework for the Training Set. The pipeline was run using SLURM batch jobs submitted to a partition on a CentOS 7 server. All resampling techniques, pre-processing, model specification, hyperparameter tuning, and evaluation metrics were implemented using the `tidymodels` suite of packages. The classifiers we used are:

- Random Forest (`rf`)
- Support Vector Machine (`svm`)
- XGBoost (`xgb`)
- Regularized Multinomial Regression (`mr`)

2.2.1 Resampling of Training Set

We used a nested cross-validation design to assess each classifier while also performing hyperparameter tuning. An outer 5-fold CV stratified by histotype was used together with an inner 5-fold CV with 2 repeats stratified by histotype. This design was chosen such that the test sets of the inner resamples would still have a reasonable number of samples belonging to the smallest minority class.

The outer resampling method cannot be the bootstrap, because the inner training and inner test sets will likely contain the same samples as a result of sampling with replacement in the outer training set. This phenomenon might result in inflated performance as some observations are used both to train and evaluate the hyperparameter tuning in the inner loop.

2.2.2 Hyperparameter Tuning

The following specifications for each classifier were used for tuning hyperparameters:

- **rf** and **xgb**: The number of trees were fixed at 500. Other hyperparameters were tuned across 10 randomly selected points in a latin hypercube design.
- **svm**: Both the cost and sigma hyperparameters were tuned across 10 randomly selected points in a latin hypercube design. We tuned the cost parameter in the range $[1, 8]$. The range for tuning the sigma parameter was obtained from the 10% and 90% quantiles of the estimation using the `kernlab::sigest()` function.
- **mr**: We generated 10 randomly selected points in a latin hypercube design for the penalty (lambda) parameter. Then, we generated 10 evenly spaced points in $[0, 1]$ for the mixture (alpha) parameter in the regularized multinomial regression model. These two sets of 10 points were crossed to generate a tuning grid of 100 points.

The hyperparameter combination that resulted in the highest average F1-score across the inner training sets was selected for each classifier to use as the model for assessing prediction performance in the outer training loop.

2.2.3 Subsampling

Here are the specifications of the subsampling methods used to handle class imbalance:

- **None**: No subsampling is performed
- **Down-sampling**: All levels except the minority class are sampled down to the same frequency as the minority class
- **Up-sampling**: All levels except the majority class are sampled up to the same frequency as the majority class
- **SMOTE**: All levels except the majority class have synthetic data generated until they have the same frequency as the majority class
- **Hybrid**: All levels except the majority class have synthetic data generated up to 50% of the frequency of the majority class, then the majority class is sampled down to the same frequency as the rest.

The figure below helps visualize how the distribution of classes changes when we apply subsampling techniques to handle class imbalance:




Figure 2.3: Visualization of Subsampling Techniques

2.2.4 Workflows

The 4 **algorithms** and 5 **subsampling** methods are crossed to create 20 different classification **workflows**. For example, the `hybrid_xgb` workflow is a classifier that first pre-processes a training set by applying a hybrid subsampling method, and then proceeds to use the XGBoost algorithm to classify ovarian histotypes.

2.3 Two-Step Algorithm

The HGSC histotype comprises of approximately 80% of cases among ovarian carcinoma patients, while the remaining 20% of cases are relatively, evenly distributed among ENOC, CCOC, LGSC, and MUC histotypes. We can implement a two-step algorithm as such:

- Step 1: use binary classification for HGSC vs. non-HGSC
- Step 2: use multinomial classification for the remaining non-HGSC classes

Let

$$\begin{aligned}
 X_k &= \text{Training data with } k \text{ classes} \\
 C_k &= \text{Class with highest } F_1 \text{ score from training } X_k \\
 W_k &= \text{Workflow associated with } C_k
 \end{aligned} \tag{2.5}$$

Figure 2.4 shows how the two-step algorithm works:

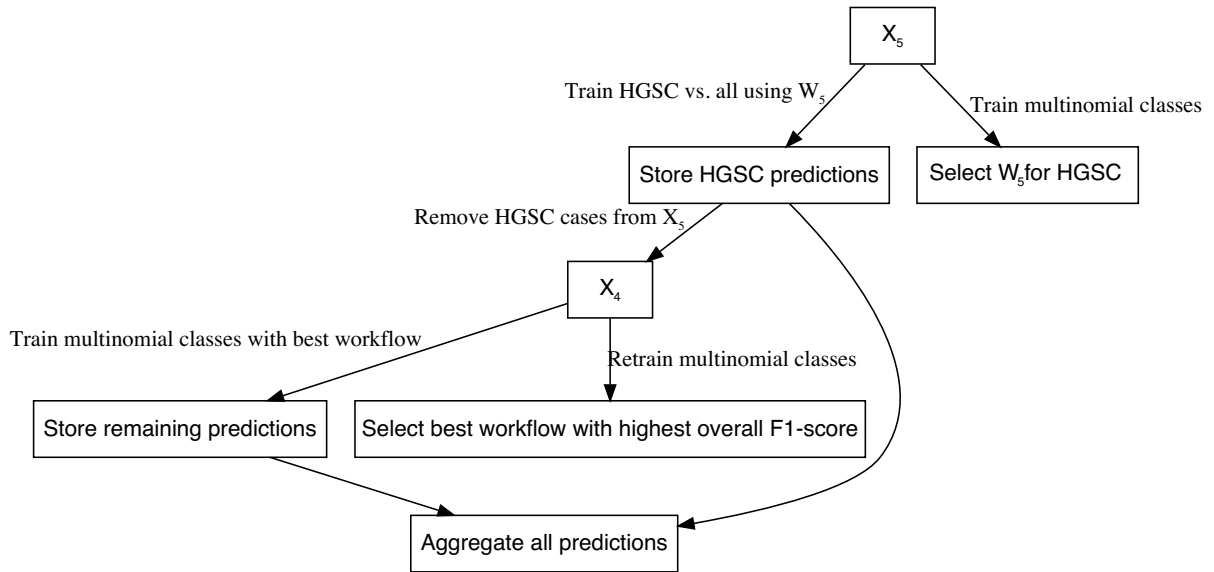


Figure 2.4: Two-Step Algorithm

2.3.1 Aggregating Predictions

The aggregation for two-step predictions is quite straightforward:

1. Predict HGSC vs. non-HGSC
2. Among all non-HGSC cases, predict CCOC vs. LGSC vs. MUC vs. ENOC



Figure 2.5: Aggregating Predictions for Two-Step Algorithm

2.4 Sequential Algorithm

Instead of training on k classes simultaneously using multinomial classifiers, we can use a sequential algorithm that performs $k-1$ one-vs-all binary classifications iteratively to obtain a final prediction of all cases. At each step in the sequence, we classify one class vs. all other classes, where the classes that make up the “other” class are those not equal to the current “one” class and excluding all “one” classes from previous steps. For example, if the “one” class in step 1 was HGSC, the “other” classes would include CCOC, ENOC, LGSC, and MUC. If the “one” class in step 2 was CCOC, the “other” classes include ENOC, LGSC, and MUC.

The order of classes and workflows to use at each step in the sequential algorithm must be determined using a retraining procedure. After removing the data associated with a particular class, we retrain using the remaining data using multinomial classifiers as described before. The class and workflow to use for the next step in the sequence is selected based on the best per-class evaluation metric value (e.g. F1-score).

Figure 2.6 illustrates how the sequential algorithm works for $K=5$, using ovarian histotypes as an example for the classes.



Figure 2.6: Sequential Algorithm

The subsampling method used in the first step of the sequential algorithm is used in all subsequent steps in order to maintain data pre-processing consistency. As a result, we are only comparing classification algorithms within one subsampling method across the entire sequential algorithm.

2.4.1 Aggregating Predictions

We have to aggregate the one-vs-all predictions from each of the sequential algorithm workflows in order to obtain a final class prediction on a holdout test set. Each sequential workflow has to be assessed on every sample to ensure that cases classified into the “all” class from a previous step of the sequence are eventually assigned a predicted class. For example, say that based on certain class-specific metrics we determined that the order of classes in the sequential algorithm was to predict HGSC vs. non-HGSC, CCOC vs. non-CCOC, LGSC vs. non-LGSC, and then MUC vs. ENOC. Figure 2.7 illustrates how the final predictions are assigned:



Figure 2.7: Aggregating Predictions for Sequential Algorithm

2.5 Performance Evaluation

2.5.1 Class Metrics

We use the accuracy, sensitivity, specificity, F1-score, kappa, balanced accuracy, and geometric mean, as class metrics to measure both training and test performance between different workflows. Multiclass extensions of these metrics can be calculated except for F1-score, where we use macro-averaging to obtain an overall metric. Class-specific metrics are calculated by recoding classes into one-vs-all categories for each class.

2.5.1.1 Accuracy

The accuracy is defined as the proportion of correct predictions out of all cases:

$$\text{accuracy} = \frac{TP}{TP + FP + FN + TN} \quad (2.6)$$

2.5.1.2 Sensitivity

Sensitivity is the proportional of correctly predicted positive cases, out of all cases that were truly positive

$$\text{sensitivity} = \frac{TP}{TP + FN} \quad (2.7)$$

2.5.1.3 Specificity

Specificity is the proportional of correctly predicted negative cases, out of all cases that were truly negative.

$$\text{specificity} = \frac{TN}{TN + FP} \quad (2.8)$$

2.5.1.4 F1-Score

The F-measure can be thought of as a harmonic mean between precision and recall:

$$F_{meas} = \frac{(1 + \beta^2) \times \text{precision} \times \text{recall}}{(\beta^2 \times \text{precision}) + \text{recall}} \quad (2.9)$$

The β value can be adjusted to place more weight upon precision or recall. The most common value is β is 1, which is also commonly known as the F1-score. A multiclass extension doesn't exist for the F1-score, so we use macro-averaging to calculate this metric when there are more than two classes. For example, with k classes, the macro-averaged F1-score is equal to:

$$F_{1_{macro}} = \frac{1}{k} \sum_{i=1}^k F_{1_i} \quad (2.10)$$

where each F_{1_i} is the F1-score computed from recoding classes into $k = i$ vs. $k \neq i$.

In situations where there is not at least one predicted case for each of the classes (e.g. for a poor classifier), F_{1_i} is undefined because the per-class precision of class i is undefined. Those F_{1_i} terms are removed from the $F_{1_{macro}}$ equation and the resulting value may be inflated. Interpreting the F1-score in such a case would be misleading.

2.5.1.5 Balanced Accuracy

Balanced accuracy is the arithmetic mean of sensitivity and specificity.

$$\text{Balanced Accuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2} \quad (2.11)$$

2.5.1.6 Kappa

Kappa is defined as:

$$\text{kappa} = \frac{p_0 - p_e}{1 - p_e} \quad (2.12)$$

where p_0 is the observed agreement among raters and p_e is the hypothetical probability of agreement due to random chance.

2.5.2 AUC

The area under the receiver operating curve (AUC) is calculated by adding up the area under the curve formed by plotting sensitivity vs. 1 - specificity. The Hand-till method is used as a multiclass extension for the AUC.

We did not use AUC to measure class-specific training set performance because combining predicted probabilities in a one-vs-all fashion might be potentially misleading. The sum of probabilities that add up to the “other” class is not equivalent to the predicted probability of the “other” class when using a multiclass classifier.

Instead, we only reported ROC curves and their associated AUCs for test set performance among the highest ranked algorithms.

2.6 Rank Aggregation

To select the best algorithm, we implemented a two-stage rank aggregation procedure using the Genetic Algorithm. First, we ranked all workflows based on per-class F1-scores, balanced accuracy, and kappa to see which workflows performed well in predicting all five histotypes. Then, we took the ranks from these three performance metrics and performed a second run of rank aggregation. The top 5 workflows were determined from the final rank aggregation result.

2.7 Gene Optimization

We want to discover an optimal set of genes for the classifiers while including specific genes from other studies such as PrOTYPE and SPOT. A total of 72 genes are used in the classifier training set.

There are 16 genes in the classifier set that overlap with the PrOTYPE classifier: COL11A1, CD74, CD2, TIMP3, LUM, CYTIP, COL3A1, THBS2, TCF7L1, HMGA2, FN1, POSTN, COL1A2, COL5A2, PDZK1IP1, FBN1.

There are also 13 genes in the classifier set that overlap with the SPOT signature: HIF1A, CXCL10, DUSP4, SOX17, MITF, CDKN3, BRCA2, CEACAM5, ANXA4, SERPINE1, TCF7L1, CRABP2, DNAJC9.

We obtain a total of 28 genes from the union of PrOTYPE and SPOT genes that we want to include in the final classifier, regardless of model performance. We then incrementally add genes one at a time from the remaining 44 candidate genes based on a variable importance rank to the set of 28 base genes and recalculate performance metrics. The number of genes at which the performance peaks or starts to plateau may indicate an optimal gene set model for us to compare with the full set model.

Here is the breakdown of genes used and whether they belong to the PrOTYPE and/or SPOT sets:

Table 2.1: Gene Distribution

Genes	PrOTYPE	SPOT
TCF7L1	v	v
COL11A1	v	
CD74	v	
CD2	v	
TIMP3	v	
LUM	v	
CYTIP	v	
COL3A1	v	
THBS2	v	
HMGA2	v	
FN1	v	
POSTN	v	
COL1A2	v	
COL5A2	v	
PDZK1IP1	v	
FBN1	v	
HIF1A		v
CXCL10		v
DUSP4		v
SOX17		v
MITF		v
CDKN3		v
BRCA2		v
CEACAM5		v
ANXA4		v
SERPINE1		v
CRABP2		v
DNAJC9		v
C10orf116		
GAD1		
TPX2		
KGFLP2		
EGFL6		
KLK7		
PBX1		

LIN28B
TFF3
MUC5B
FUT3
STC1
BCL2
PAX8
GCNT3
GPR64
ADCYAP1R1
IGKC
BRCA1
IGJ
TFF1
MET
CYP2C18
CYP4B1
SLC3A1
EPAS1
HNF1B
IL6
ATP5G3
DKK4
SENP8
CAPN2
C1orf173
CPNE8
IGFBP1
WT1
TP53
SEMA6A
SERPINA5
ZBED1
TSPAN8
SCGB1D2
LGALS4
MAP1LC3A

2.7.1 Variable Importance

Variable importance is calculated using either a model-based approach if it is available, or a permutation-based VI score otherwise. The variable importance scores are averaged across the outer training folds, and then ranked from highest to lowest.

For the sequential and two-step classifiers, we calculate an overall VI rank by taking the cumulative union of genes at each variable importance rank across all sequences, until all genes have been included.

The variable importance measures are:

- Random Forest: impurity measure (Gini index)
- XGBoost: gain (fractional contribution of each feature to the model based on the total gain of the corresponding features's splits)
- SVM: permutation based p-values
- Multinomial regression: absolute value of estimated coefficients at cross-validated lambda value

3 Distributions

3.1 Histotype Distribution

Table 3.1: Histotype Distribution in Training Set by Processing Stage

Variable	Levels	CS1	CS2	CS3	Total
Selected Cohorts					
Histotype	HGSC	126 (43%)	655 (73%)	1779 (72%)	2560 (70%)
	CCOC	48 (16%)	61 (7%)	181 (7%)	290 (8%)
	ENOC	60 (20%)	34 (4%)	268 (11%)	362 (10%)
	MUC	20 (7%)	62 (7%)	77 (3%)	159 (4%)
	LGSC	21 (7%)	21 (2%)	42 (2%)	84 (2%)
	Other	19 (6%)	70 (8%)	130 (5%)	219 (6%)
Total	N (%)	294 (8%)	903 (25%)	2477 (67%)	3674 (100%)
QC					
Histotype	HGSC	120 (42%)	641 (73%)	1636 (72%)	2397 (70%)
	CCOC	48 (17%)	61 (7%)	173 (8%)	282 (8%)
	ENOC	60 (21%)	32 (4%)	229 (10%)	321 (9%)
	MUC	19 (7%)	60 (7%)	69 (3%)	148 (4%)
	LGSC	20 (7%)	21 (2%)	40 (2%)	81 (2%)
	Other	19 (7%)	67 (8%)	126 (6%)	212 (6%)
Total	N (%)	286 (8%)	882 (26%)	2273 (66%)	3441 (100%)
Main Histotypes					
Histotype	HGSC	120 (45%)	641 (79%)	1636 (76%)	2397 (74%)
	CCOC	48 (18%)	61 (7%)	173 (8%)	282 (9%)
	ENOC	60 (22%)	32 (4%)	229 (11%)	321 (10%)
	MUC	19 (7%)	60 (7%)	69 (3%)	148 (5%)
	LGSC	20 (7%)	21 (3%)	40 (2%)	81 (3%)
Total	N (%)	267 (8%)	815 (25%)	2147 (66%)	3229 (100%)
Removed Duplicates					
	HGSC	117 (47%)	623 (79%)	1540 (77%)	2280 (75%)

Histotype	CCOC	45 (18%)	55 (7%)	159 (8%)	259 (9%)
	ENOC	56 (22%)	28 (4%)	216 (11%)	300 (10%)
	MUC	16 (6%)	58 (7%)	59 (3%)	133 (4%)
	LGSC	15 (6%)	20 (3%)	36 (2%)	71 (2%)
Total	N (%)	249 (8%)	784 (26%)	2010 (66%)	3043 (100%)
Normalized and Recombined					
Histotype	HGSC	116 (48%)	622 (80%)	451 (96%)	1189 (80%)
	CCOC	44 (18%)	54 (7%)	4 (1%)	102 (7%)
	ENOC	55 (23%)	27 (3%)	4 (1%)	86 (6%)
	MUC	15 (6%)	57 (7%)	5 (1%)	77 (5%)
	LGSC	14 (6%)	19 (2%)	4 (1%)	37 (2%)
Total	N (%)	244 (16%)	779 (52%)	468 (31%)	1491 (100%)
Removed Replicates					
Histotype	HGSC	9 (12%)	552 (79%)	451 (96%)	1012 (81%)
	CCOC	25 (32%)	52 (7%)	4 (1%)	81 (7%)
	ENOC	37 (48%)	25 (4%)	4 (1%)	66 (5%)
	MUC	3 (4%)	53 (8%)	5 (1%)	61 (5%)
	LGSC	3 (4%)	16 (2%)	4 (1%)	23 (2%)
Total	N (%)	77 (6%)	698 (56%)	468 (38%)	1243 (100%)

Table 3.2: Histotype Distribution in Training, Confirmation, and Validation Sets

Variable	Levels	Training	Confirmation	Validation
Histotype	HGSC	1012 (81%)	422 (66%)	666 (74%)
	CCOC	81 (7%)	75 (12%)	79 (9%)
	ENOC	66 (5%)	106 (16%)	105 (12%)
	MUC	61 (5%)	27 (4%)	26 (3%)
	LGSC	23 (2%)	13 (2%)	18 (2%)
Total	N (%)	1243 (45%)	643 (23%)	894 (32%)

3.2 Cohort Distribution

Table 3.3: Pre-QC Cohort Distribution by CodeSet

CodeSet	CS1, N = 294	CS2, N = 903	CS3, N = 2,477
Cohort			
OOU	108 (37%)	43 (4.8%)	19 (0.8%)
OOUE	32 (11%)	30 (3.3%)	11 (0.4%)
VOA	145 (49%)	122 (14%)	538 (22%)
OVAR3	0 (0%)	150 (17%)	0 (0%)
ICON7	0 (0%)	416 (46%)	0 (0%)
MAYO	6 (2.0%)	63 (7.0%)	0 (0%)
DOVE4	0 (0%)	0 (0%)	1,160 (47%)
TNCO	0 (0%)	0 (0%)	691 (28%)
MTL	3 (1.0%)	59 (6.5%)	0 (0%)
JAPAN	0 (0%)	8 (0.9%)	0 (0%)
POOL-CTRL	0 (0%)	12 (1.3%)	0 (0%)
POOL-1	0 (0%)	0 (0%)	31 (1.3%)
POOL-2	0 (0%)	0 (0%)	14 (0.6%)
POOL-3	0 (0%)	0 (0%)	13 (0.5%)

¹ n (%)

3.3 Quality Control

3.3.1 Failed Samples

We use an aggregated `QCFlag` that considers a sample to have failed QC if any of the following QC conditions are flagged:

- Linearity
- Imaging
- Smallest Positive Control
- Normality

Table 3.4: Quality Control Summary

Quality Control Flag	CS1, N = 294	CS2, N = 903	CS3, N = 2,477
Linearity			
Failed	0 (0%)	4 (0.4%)	0 (0%)
Passed	294 (100%)	899 (100%)	2,477 (100%)
Imaging			
Failed	3 (1.0%)	0 (0%)	4 (0.2%)
Passed	291 (99%)	903 (100%)	2,473 (100%)
Smallest Positive Control			
Failed	0 (0%)	2 (0.2%)	0 (0%)
Passed	294 (100%)	901 (100%)	2,477 (100%)
Normality			
Failed	5 (1.7%)	19 (2.1%)	200 (8.1%)
Passed	289 (98%)	884 (98%)	2,277 (92%)
Overall QC			
Failed	8 (2.7%)	21 (2.3%)	204 (8.2%)
Passed	286 (97%)	882 (98%)	2,273 (92%)

¹ n (%)

3.3.2 %GD vs. SNR

% Genes Detected vs. Signal-to-Noise Ratio

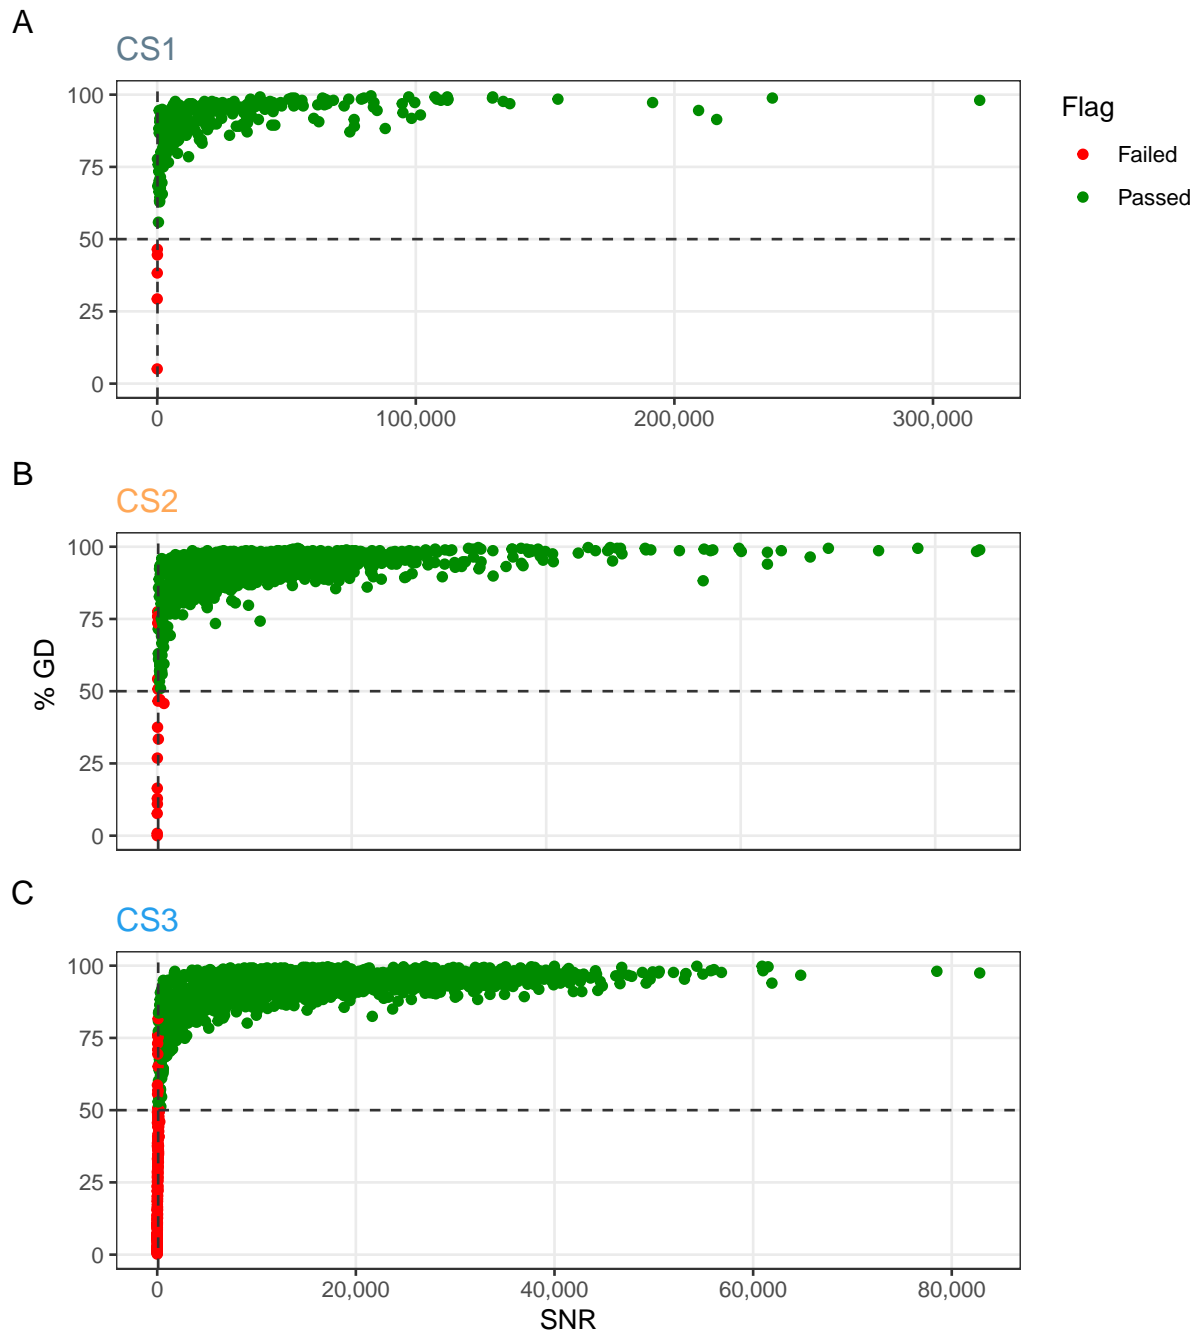


Figure 3.1: % Genes Detected vs. Signal to Noise Ratio

% Genes Detected vs. Signal-to-Noise Ratio (Zoomed)

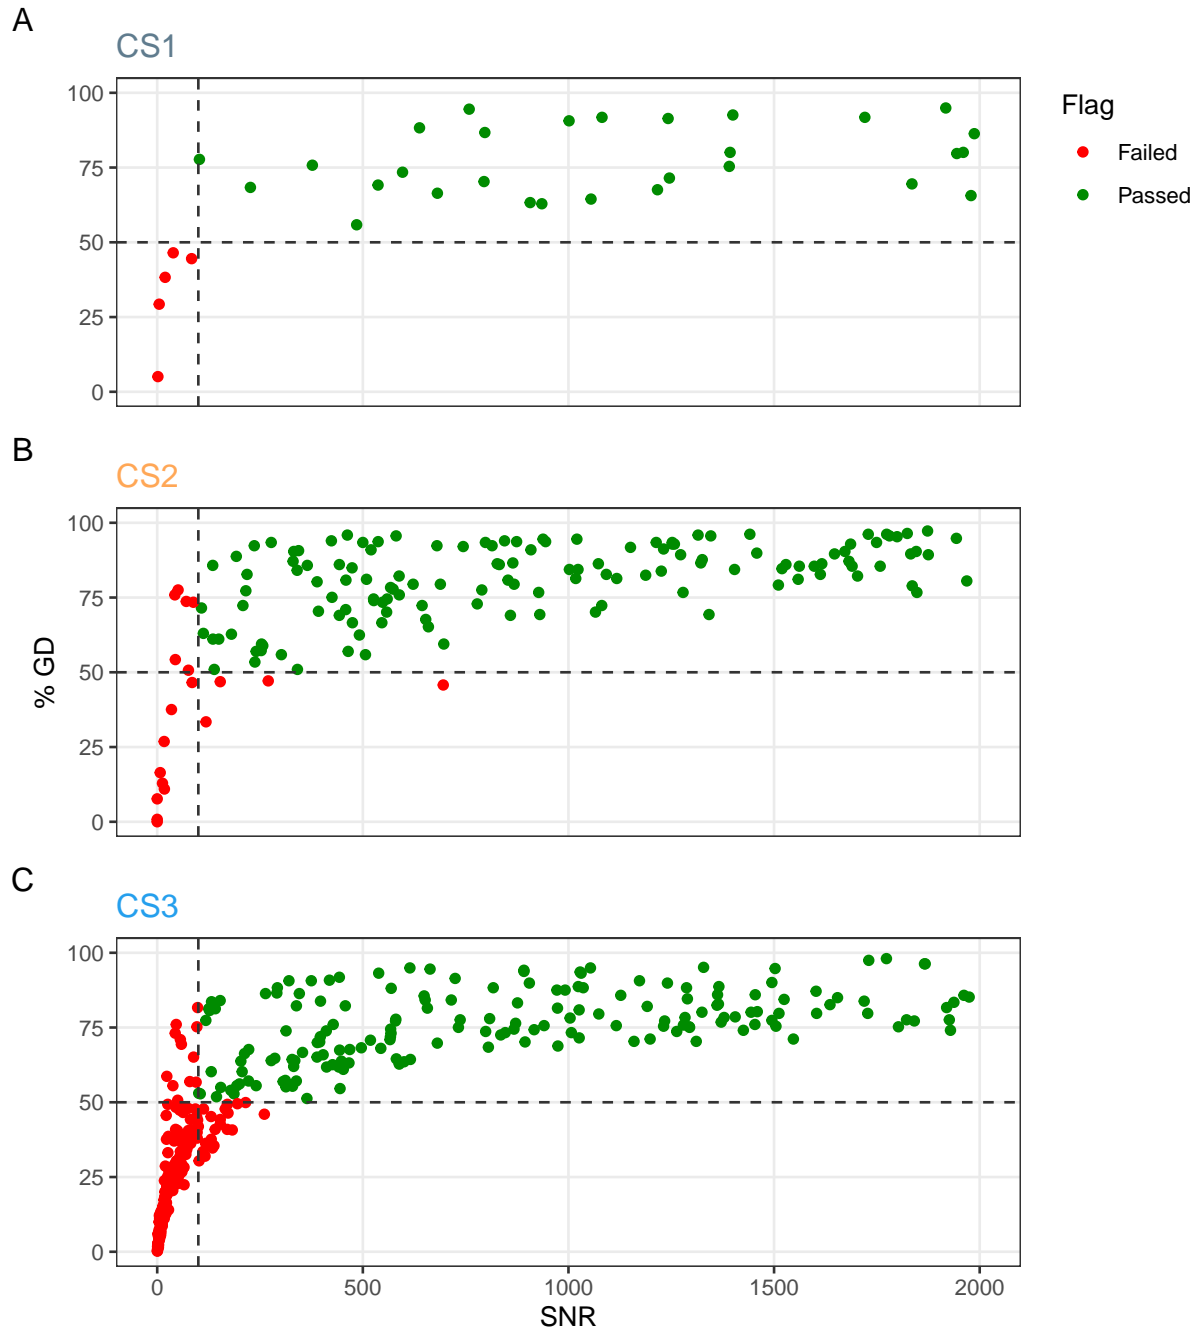


Figure 3.2: % Genes Detected vs. Signal to Noise Ratio (Zoomed)

3.4 Pairwise Gene Expression

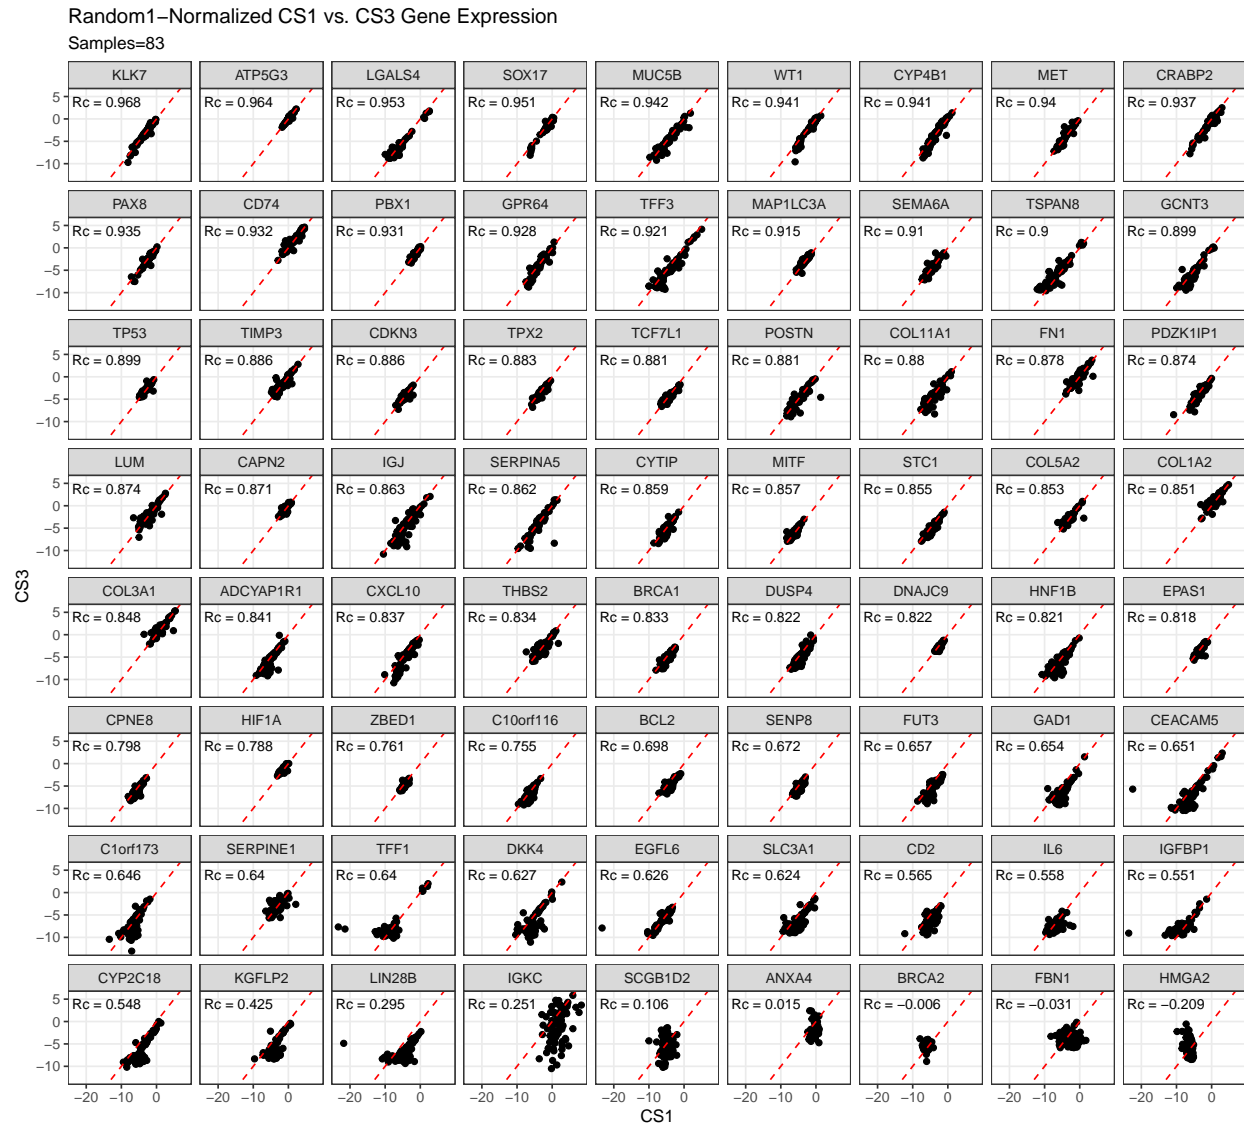


Figure 3.3: Random1-Normalized CS1 vs. CS3 Gene Expression

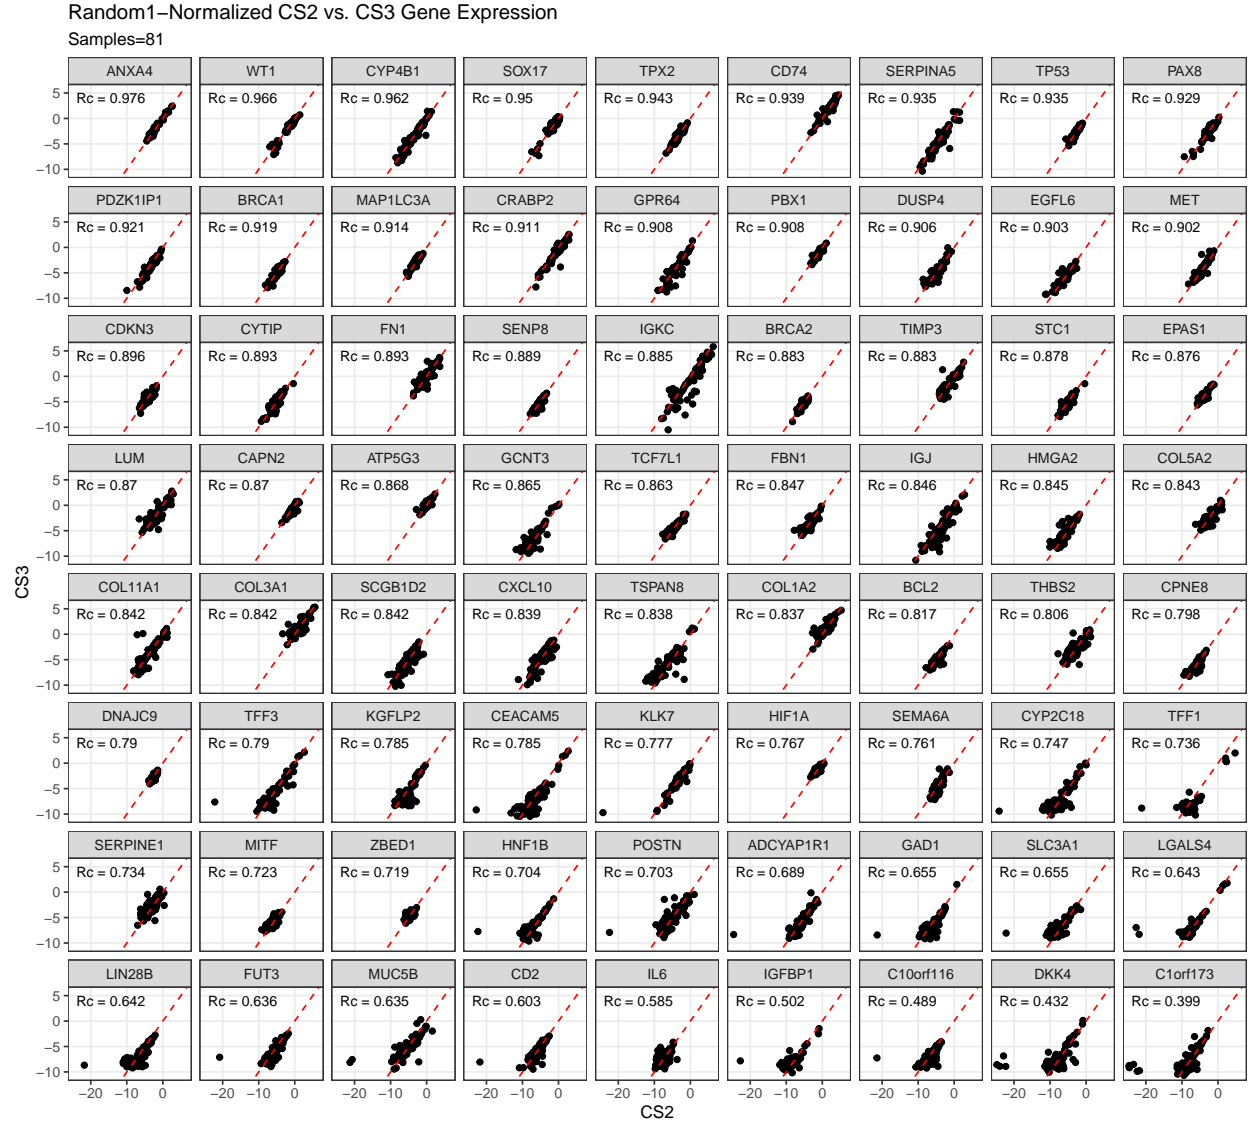


Figure 3.4: Random1-Normalized CS2 vs. CS3 Gene Expression

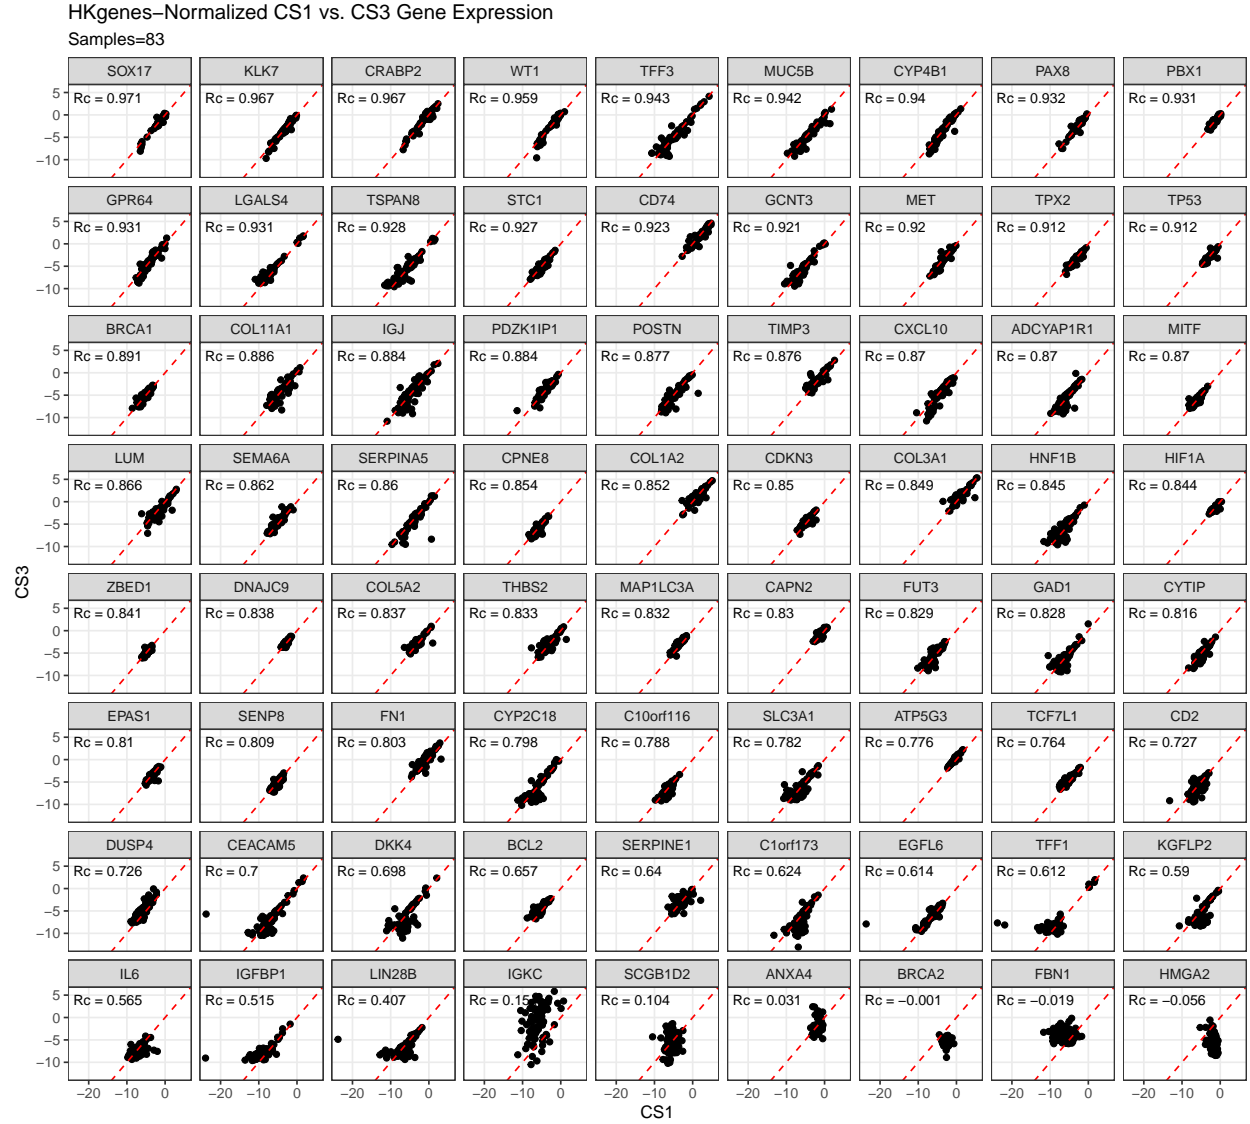


Figure 3.5: HKgenes-Normalized CS1 vs. CS3 Gene Expression



Figure 3.6: HKgenes-Normalized CS2 vs. CS3 Gene Expression

4 Results

We summarize cross-validated training performance of class metrics in the training set. The accuracy, F1-score, and kappa, are the metrics of interest. Workflows are ordered by their mean estimates across the outer folds of the nested CV for each metric.

4.1 Training Set

4.1.1 Accuracy

Table 4.1: Training Set Mean Accuracy

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.93	0.947	0.982	0.965	0.981	0.985
	svm	0.941	0.961	0.982	0.97	0.985	0.985
	xgb	0.823	0.824	0.943	0.947	0.982	0.951
	mr	0.814	0.814	0.935	0.947	0.982	0.951
down	rf	0.829	0.861	0.976	0.945	0.909	0.968
	svm	0.797	0.831	0.965	0.924	0.901	0.973
	xgb	0.206	0.32	0.583	0.947	0.79	0.773
	mr	0.817	0.851	0.969	0.928	0.92	0.966
up	rf	0.928	0.949	0.979	0.969	0.977	0.982
	svm	0.932	0.959	0.973	0.963	0.986	0.982
	xgb	0.937	0.962	0.982	0.967	0.982	0.982
	mr	0.889	0.92	0.974	0.95	0.961	0.972
smote	rf	0.94	0.963	0.981	0.97	0.981	0.985
	svm	0.932	0.96	0.974	0.962	0.986	0.982
	xgb	0.934	0.957	0.981	0.967	0.982	0.981
	mr	0.885	0.914	0.974	0.956	0.947	0.979
hybrid	rf	0.932	0.957	0.979	0.969	0.981	0.978
	svm	0.92	0.948	0.973	0.957	0.981	0.982
	xgb	0.928	0.953	0.978	0.966	0.978	0.98
	mr	0.885	0.913	0.978	0.954	0.949	0.976




Figure 4.1: Training Set Mean Accuracy

4.1.2 Sensitivity

Table 4.2: Training Set Mean Sensitivity

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.666	0.991	0.797	0.585	0.133	0.822
	svm	0.713	0.993	0.802	0.695	0.264	0.81
	xgb	0.241	1	0.203	0	0	0
	mr	0.2	1	0	0	0	0
down	rf	0.805	0.838	0.816	0.684	0.865	0.824
	svm	0.814	0.8	0.743	0.777	0.971	0.781
	xgb	0.2	0.2	0.4	0	0.2	0.2
	mr	0.79	0.826	0.772	0.683	0.865	0.804
up	rf	0.695	0.981	0.805	0.636	0.28	0.77
	svm	0.725	0.989	0.748	0.617	0.528	0.744
	xgb	0.75	0.98	0.805	0.715	0.397	0.854
	mr	0.834	0.91	0.789	0.766	0.871	0.833
smote	rf	0.755	0.984	0.805	0.762	0.38	0.841
	svm	0.738	0.988	0.748	0.661	0.534	0.76
	xgb	0.796	0.965	0.84	0.738	0.596	0.839
	mr	0.806	0.905	0.792	0.79	0.703	0.839
hybrid	rf	0.775	0.97	0.801	0.777	0.47	0.855
	svm	0.811	0.957	0.783	0.764	0.763	0.788
	xgb	0.792	0.96	0.825	0.727	0.594	0.854
	mr	0.821	0.903	0.812	0.782	0.77	0.839

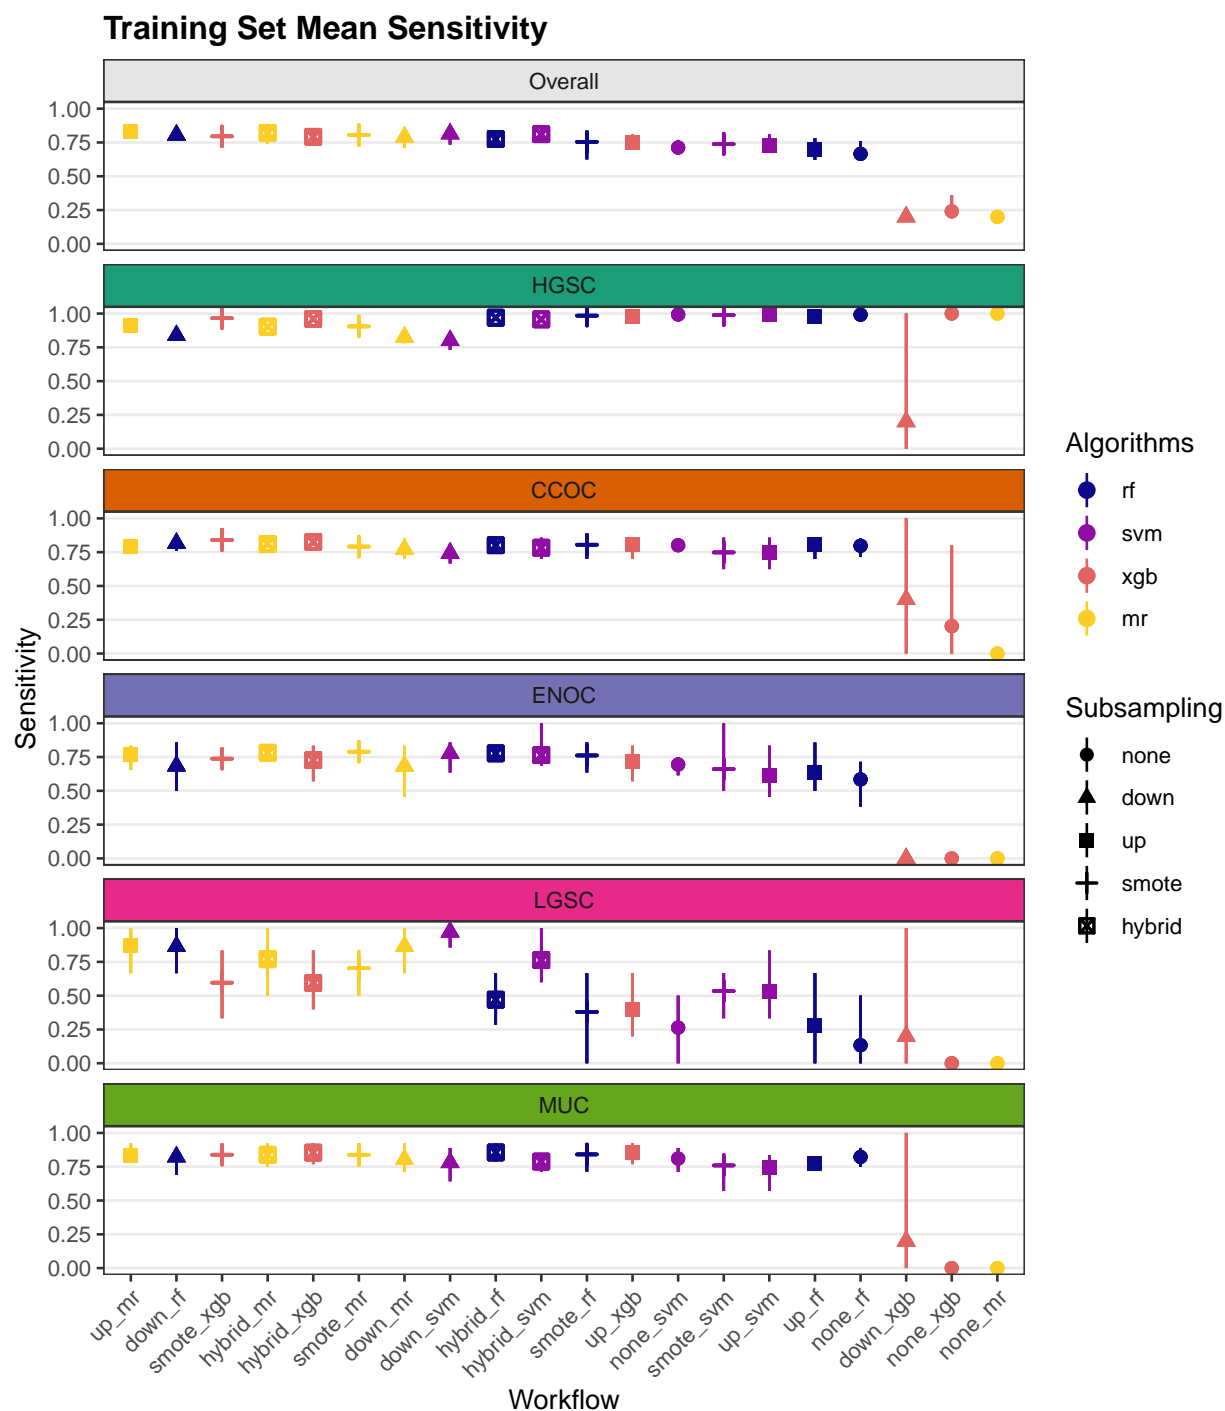


Figure 4.2: Training Set Mean Sensitivity

4.1.3 Specificity

Table 4.3: Training Set Mean Specificity

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.947	0.761	0.995	0.988	0.998	0.993
	svm	0.959	0.821	0.995	0.986	0.998	0.994
	xgb	0.811	0.055	0.999	1	1	1
	mr	0.8	0	1	1	1	1
down	rf	0.958	0.961	0.987	0.958	0.91	0.975
	svm	0.952	0.963	0.98	0.933	0.9	0.984
	xgb	0.8	0.8	0.6	1	0.8	0.8
	mr	0.955	0.959	0.983	0.94	0.92	0.975
up	rf	0.953	0.803	0.991	0.987	0.991	0.992
	svm	0.959	0.833	0.99	0.984	0.993	0.995
	xgb	0.968	0.881	0.995	0.981	0.993	0.988
	mr	0.971	0.964	0.987	0.962	0.963	0.979
smote	rf	0.966	0.866	0.993	0.983	0.993	0.992
	svm	0.96	0.841	0.991	0.982	0.994	0.994
	xgb	0.974	0.923	0.991	0.979	0.989	0.988
	mr	0.968	0.95	0.987	0.965	0.951	0.986
hybrid	rf	0.97	0.903	0.991	0.981	0.991	0.985
	svm	0.969	0.91	0.986	0.97	0.985	0.992
	xgb	0.971	0.916	0.989	0.979	0.985	0.987
	mr	0.968	0.954	0.99	0.964	0.952	0.983




Figure 4.3: Training Set Mean Specificity

4.1.4 F1-Score

Table 4.4: Training Set Mean F1-Score

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.751	0.968	0.849	0.622	0.262	0.838
	svm	0.775	0.976	0.852	0.695	0.433	0.837
	xgb	0.841	0.902	0.611	NaN	NaN	NaN
	mr	0.897	0.897	NaN	NaN	NaN	NaN
down	rf	0.648	0.907	0.813	0.545	0.255	0.719
	svm	0.623	0.884	0.728	0.486	0.275	0.744
	xgb	0.257	0.91	0.114	NaN	0.04	0.106
	mr	0.622	0.9	0.76	0.474	0.273	0.704
up	rf	0.707	0.969	0.829	0.657	0.278	0.803
	svm	0.741	0.975	0.776	0.612	0.54	0.802
	xgb	0.747	0.977	0.855	0.678	0.408	0.817
	mr	0.706	0.949	0.804	0.587	0.448	0.742
smote	rf	0.748	0.978	0.837	0.716	0.37	0.84
	svm	0.751	0.975	0.783	0.626	0.566	0.805
	xgb	0.763	0.973	0.845	0.663	0.528	0.806
	mr	0.7	0.945	0.797	0.631	0.334	0.793
hybrid	rf	0.748	0.974	0.828	0.698	0.448	0.791
	svm	0.751	0.967	0.785	0.622	0.569	0.81
	xgb	0.753	0.97	0.829	0.678	0.484	0.803
	mr	0.708	0.944	0.829	0.628	0.367	0.769

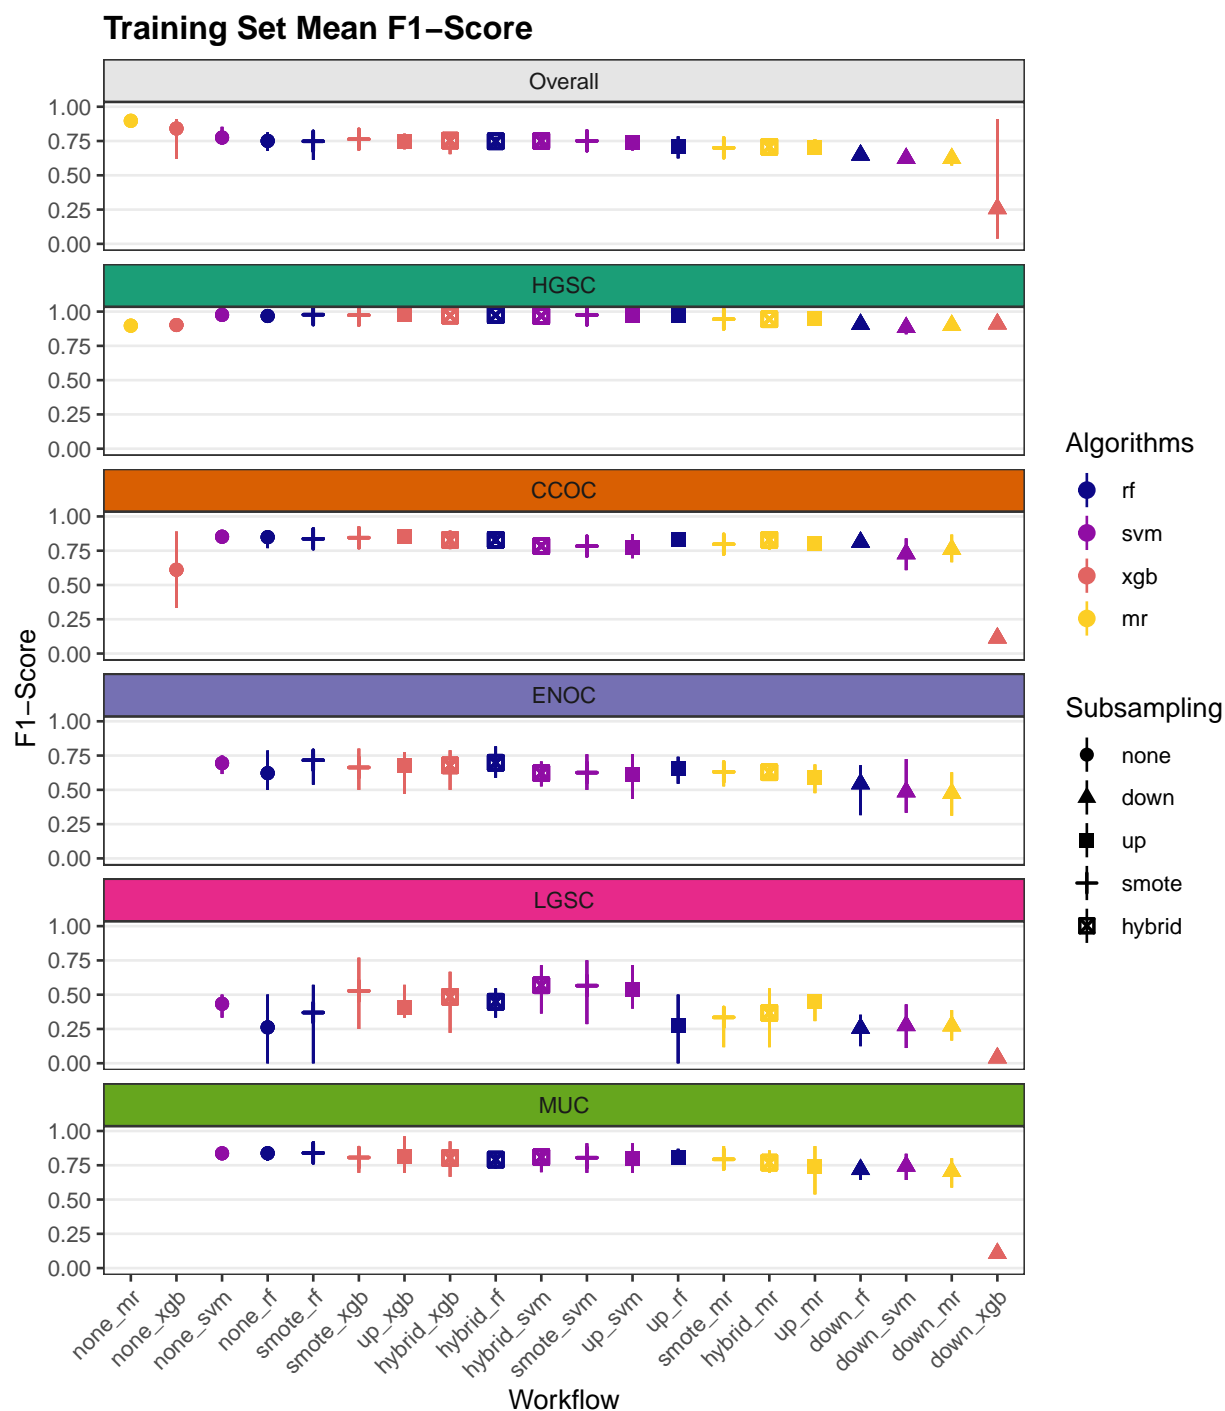


Figure 4.4: Training Set Mean F1-Score

4.1.5 Balanced Accuracy

Table 4.5: Training Set Mean Balanced Accuracy

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.806	0.876	0.896	0.786	0.566	0.908
	svm	0.836	0.907	0.898	0.841	0.631	0.902
	xgb	0.526	0.528	0.601	0.5	0.5	0.5
	mr	0.5	0.5	0.5	0.5	0.5	0.5
down	rf	0.882	0.899	0.902	0.821	0.887	0.9
	svm	0.883	0.882	0.862	0.855	0.936	0.882
	xgb	0.5	0.5	0.5	0.5	0.5	0.5
	mr	0.873	0.892	0.878	0.812	0.893	0.889
up	rf	0.824	0.892	0.898	0.812	0.636	0.881
	svm	0.842	0.911	0.869	0.801	0.761	0.87
	xgb	0.859	0.931	0.9	0.848	0.695	0.921
	mr	0.903	0.937	0.888	0.864	0.917	0.906
smote	rf	0.86	0.925	0.899	0.872	0.687	0.917
	svm	0.849	0.915	0.869	0.822	0.764	0.877
	xgb	0.885	0.944	0.915	0.858	0.792	0.914
	mr	0.887	0.927	0.889	0.877	0.827	0.913
hybrid	rf	0.872	0.937	0.896	0.879	0.731	0.92
	svm	0.89	0.933	0.885	0.867	0.874	0.89
	xgb	0.882	0.938	0.907	0.853	0.79	0.92
	mr	0.895	0.928	0.901	0.873	0.861	0.911

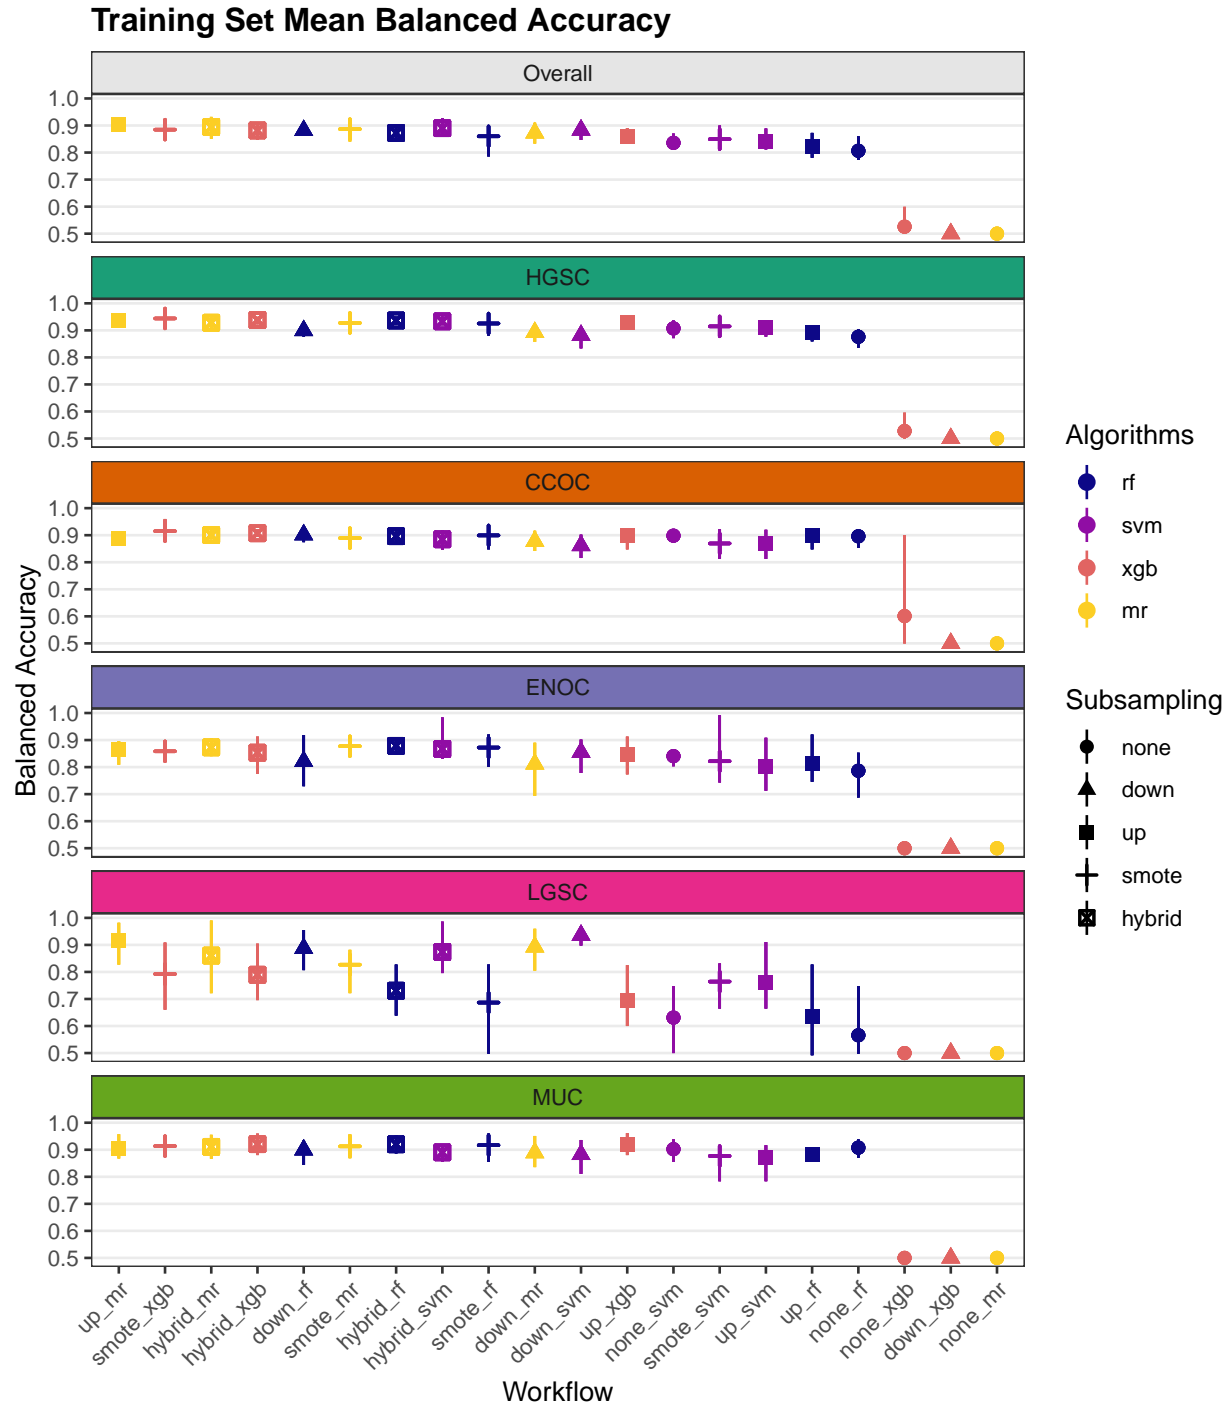


Figure 4.5: Training Set Mean Balanced Accuracy

4.1.6 Kappa

Table 4.6: Training Set Mean Kappa

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.767	0.811	0.84	0.604	0.154	0.83
	svm	0.808	0.862	0.843	0.679	0.342	0.829
	xgb	0.084	0.083	0.24	0	0	0
	mr	0	0	0	0	0	0
down	rf	0.597	0.633	0.8	0.517	0.232	0.702
	svm	0.547	0.577	0.709	0.452	0.252	0.73
	xgb	0	0	0	0	0	0
	mr	0.574	0.613	0.744	0.439	0.251	0.687
up	rf	0.764	0.819	0.818	0.641	0.269	0.793
	svm	0.779	0.859	0.762	0.593	0.533	0.793
	xgb	0.804	0.871	0.845	0.66	0.399	0.807
	mr	0.706	0.766	0.79	0.562	0.434	0.728
smote	rf	0.806	0.87	0.827	0.7	0.363	0.832
	svm	0.783	0.862	0.77	0.606	0.56	0.796
	xgb	0.803	0.862	0.835	0.646	0.519	0.796
	mr	0.695	0.747	0.783	0.608	0.316	0.782
hybrid	rf	0.793	0.858	0.817	0.682	0.439	0.78
	svm	0.765	0.831	0.77	0.601	0.56	0.801
	xgb	0.786	0.846	0.817	0.661	0.475	0.793
	mr	0.696	0.745	0.818	0.604	0.349	0.757




Figure 4.6: Training Set Mean Kappa

4.1.7 G-mean

DEPRECATED

Table 4.7: Training Set Mean G-mean

Subsampling	Algorithms	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
none	rf	0.261	0.868	0.89	0.755	0.223	0.903
	svm	0.54	0.903	0.893	0.827	0.453	0.897
	xgb	0	0.146	0.271	0	0	0
	mr	0	0	0	0	0	0
down	rf	0.798	0.897	0.897	0.806	0.884	0.895
	svm	0.809	0.878	0.853	0.85	0.934	0.875
	xgb	0	0	0	0	0	0
	mr	0.782	0.89	0.871	0.796	0.889	0.884
up	rf	0.421	0.887	0.893	0.789	0.403	0.874
	svm	0.699	0.907	0.859	0.775	0.715	0.859
	xgb	0.706	0.929	0.894	0.836	0.614	0.918
	mr	0.829	0.936	0.882	0.858	0.913	0.903
smote	rf	0.607	0.923	0.894	0.864	0.545	0.913
	svm	0.714	0.911	0.859	0.799	0.724	0.867
	xgb	0.778	0.944	0.912	0.849	0.759	0.91
	mr	0.801	0.927	0.884	0.872	0.814	0.909
hybrid	rf	0.749	0.936	0.891	0.872	0.676	0.917
	svm	0.804	0.933	0.878	0.859	0.863	0.884
	xgb	0.778	0.938	0.903	0.842	0.759	0.918
	mr	0.816	0.928	0.896	0.868	0.851	0.907




Figure 4.7: Training Set Mean G-mean

4.2 Rank Aggregation

Multi-step methods:

- **sequential:** sequential algorithm sequence of subsampling methods and algorithms used are:
 - HGSC vs. non-HGSC using SMOTE subsampling and random forest
 - CCOC vs. non-CCOC using hybrid subsampling and XGBoost
 - ENOC vs. non-ENOC using upsampling and support vector machine
 - LGSC vs. MUC using hybrid subsampling and regularized multinomial regression
- **two_step:** two-step algorithm sequence of subsampling methods and algorithms used are:
 - HGSC vs. non-HGSC using SMOTE subsampling and random forest
 - CCOC vs. ENOC vs. MUC vs. LGSC using hybrid subsampling and support vector machine

We conduct rank aggregation using a two-stage nested approach:

1. First we rank aggregate the per-class metrics for F1-score, balanced accuracy and kappa.
2. Then we take the aggregated lists from the three metrics and perform a final rank aggregation.
3. The top workflows from the final rank aggregation are used for gene optimization in the confirmation set

4.2.1 Across Classes

4.2.1.1 F1-Score

Table 4.8: F1-Score Rank Aggregation Summary

Workflow	Rank	HGSC	CCOC	ENOC	LGSC	MUC
sequential	1	0.969	0.855	0.891	0.919	0.967
two_step	2	0.969	0.865	0.738	0.782	0.864
smote_rf	3	0.978	0.837	0.716	0.37	0.84
none_svm	4	0.976	0.852	0.695	0.433	0.837
up_xgb	5	0.977	0.855	0.678	0.408	0.817
smote_xgb	6	0.973	0.845	0.663	0.528	0.806
hybrid_xgb	7	0.97	0.829	0.678	0.484	0.803
smote_svm	8	0.975	0.783	0.626	0.566	0.805
hybrid_rf	9	0.974	0.828	0.698	0.448	0.791
up_rf	10	0.969	0.829	0.657	0.278	0.803
up_svm	11	0.975	0.776	0.612	0.54	0.802
hybrid_mr	12	0.944	0.829	0.628	0.367	0.769
hybrid_svm	13	0.967	0.785	0.622	0.569	0.81
none_rf	14	0.968	0.849	0.622	0.262	0.838
smote_mr	15	0.945	0.797	0.631	0.334	0.793
up_mr	16	0.949	0.804	0.587	0.448	0.742
down_rf	17	0.907	0.813	0.545	0.255	0.719
down_mr	18	0.9	0.76	0.474	0.273	0.704
down_svm	19	0.884	0.728	0.486	0.275	0.744

4.2.1.2 Balanced Accuracy

Table 4.9: Balanced Accuracy Rank Aggregation Summary

Workflow	Rank	HGSC	CCOC	ENOC	LGSC	MUC
All	All	All	All	All	All	All
sequential	1	0.919	0.889	0.905	0.955	0.955
hybrid_xgb	2	0.938	0.907	0.853	0.79	0.92
smote_xgb	3	0.944	0.915	0.858	0.792	0.914
hybrid_rf	4	0.937	0.896	0.879	0.731	0.92
smote_rf	5	0.925	0.899	0.872	0.687	0.917
up_xgb	6	0.931	0.9	0.848	0.695	0.921
up_mr	7	0.937	0.888	0.864	0.917	0.906
hybrid_mr	8	0.928	0.901	0.873	0.861	0.911
smote_mr	9	0.927	0.889	0.877	0.827	0.913
two_step	10	0.919	0.893	0.819	0.924	0.908
hybrid_svm	11	0.933	0.885	0.867	0.874	0.89
none_svm	12	0.907	0.898	0.841	0.631	0.902
smote_svm	13	0.915	0.869	0.822	0.764	0.877
down_rf	14	0.899	0.902	0.821	0.887	0.9
down_mr	15	0.892	0.878	0.812	0.893	0.889
down_svm	16	0.882	0.862	0.855	0.936	0.882
up_rf	17	0.892	0.898	0.812	0.636	0.881
up_svm	18	0.911	0.869	0.801	0.761	0.87
none_rf	19	0.876	0.896	0.786	0.566	0.908
down_xgb	20	0.5	0.5	0.5	0.5	0.5
none_mr	21	0.5	0.5	0.5	0.5	0.5
none_xgb	22	0.528	0.601	0.5	0.5	0.5

4.2.1.3 Kappa

Table 4.10: Kappa Rank Aggregation Summary

Workflow	Rank	HGSC	CCOC	ENOC	LGSC	MUC
All	All	All	All	All	All	All
sequential	1	0.833	0.774	0.819	0.886	0.886
smote_rf	2	0.87	0.827	0.7	0.363	0.832
up_xgb	3	0.871	0.845	0.66	0.399	0.807
none_svm	4	0.862	0.843	0.679	0.342	0.829
two_step	5	0.833	0.796	0.632	0.758	0.818
smote_xgb	6	0.862	0.835	0.646	0.519	0.796
hybrid_rf	7	0.858	0.817	0.682	0.439	0.78
hybrid_xgb	8	0.846	0.817	0.661	0.475	0.793
smote_svm	9	0.862	0.77	0.606	0.56	0.796
up_svm	10	0.859	0.762	0.593	0.533	0.793
hybrid_svm	11	0.831	0.77	0.601	0.56	0.801
up_rf	12	0.819	0.818	0.641	0.269	0.793
none_rf	13	0.811	0.84	0.604	0.154	0.83
hybrid_mr	14	0.745	0.818	0.604	0.349	0.757
smote_mr	15	0.747	0.783	0.608	0.316	0.782
up_mr	16	0.766	0.79	0.562	0.434	0.728
down_rf	17	0.633	0.8	0.517	0.232	0.702
down_mr	18	0.613	0.744	0.439	0.251	0.687
down_svm	19	0.577	0.709	0.452	0.252	0.73
down_xgb	20	0	0	0	0	0
none_mr	21	0	0	0	0	0
none_xgb	22	0.083	0.24	0	0	0

4.2.2 Across Metrics

Table 4.11: Rank Aggregation Comparison of Metrics Used

Rank	F1	Balanced Accuracy	Kappa
1	sequential	sequential	sequential
2	two_step	hybrid_xgb	smote_rf
3	smote_rf	smote_xgb	up_xgb
4	none_svm	hybrid_rf	none_svm
5	up_xgb	smote_rf	two_step
6	smote_xgb	up_xgb	smote_xgb
7	hybrid_xgb	up_mr	hybrid_rf
8	smote_svm	hybrid_mr	hybrid_xgb
9	hybrid_rf	smote_mr	smote_svm
10	up_rf	two_step	up_svm
11	up_svm	hybrid_svm	hybrid_svm
12	hybrid_mr	none_svm	up_rf
13	hybrid_svm	smote_svm	none_rf
14	none_rf	down_rf	hybrid_mr
15	smote_mr	down_mr	smote_mr
16	up_mr	down_svm	up_mr
17	down_rf	up_rf	down_rf
18	down_mr	up_svm	down_mr
19	down_svm	none_rf	down_svm
20	NA	down_xgb	down_xgb
21	NA	none_mr	none_mr
22	NA	none_xgb	none_xgb

Table 4.12: Top 5 Workflows from Final Rank Aggregation

Rank	Workflow
1	sequential
2	smote_rf
3	two_step
4	none_svm
5	up_xgb

4.2.3 Top Workflows

We look at the per-class evaluation metrics of the top 5 workflows.

Top 5 Workflow Per-Class Evaluation Metrics by Metric




Figure 4.8: Top 5 Workflow Per-Class Evaluation Metrics by Metric




Figure 4.9: Top 5 Workflow Per-Class Evaluation Metrics by Metric

Misclassified cases from a previous step of the sequence of classifiers are not included in subsequent steps of the training set CV folds. Thus, we cannot piece together the test set predictions from the sequential and two-step algorithms to obtain overall metrics.

4.3 Optimal Gene Sets

4.3.1 Sequential Algorithm




Figure 4.10: Gene Optimization for Sequential Classifier

In the sequential algorithm, sequences 1, 2, and 4 have relatively flat average F1-scores across the number of genes added. However, we can observe in sequence 3, the F1-score stabilizes at around 0.88 when we reach 7 genes added, hence the optimal number of genes used will be $n=28+7=35$. The added genes are: CYP2C18, TFF3, TP53, HNF1B, WT1, MAP1LC3A and SLC3A1.

Table 4.13: Gene Profile of Optimal Set in Sequential Algorithm

Set	Genes	PrOTYPE	SPOT	Optimal Set	Candidate Rank
	COL11A1	v		(*)	
	CD74	v		(*)	
	CD2	v		(*)	
	TIMP3	v		(*)	
	LUM	v		(*)	

Base	CYTIP	v	(*)	
	COL3A1	v	(*)	
	THBS2	v	(*)	
	TCF7L1	v	v	(*)
	HMGA2	v		(*)
	FN1	v		(*)
	POSTN	v		(*)
	COL1A2	v		(*)
	COL5A2	v		(*)
	PDZK1IP1	v		(*)
	FBN1	v		(*)
	HIF1A		v	(*)
	CXCL10		v	(*)
	DUSP4		v	(*)
	SOX17		v	(*)
	MITF		v	(*)
	CDKN3		v	(*)
	BRCA2		v	(*)
	CEACAM5		v	(*)
	ANXA4		v	(*)
	SERPINE1		v	(*)
	CRABP2		v	(*)
	DNAJC9		v	(*)
	CYP2C18			(*) 1
	TFF3			(*) 2
	TP53			(*) 3
	HNF1B			(*) 4
	WT1			(*) 5
	MAP1LC3A			(*) 6
	SLC3A1			(*) 7
	EPAS1			(*) 8
	EGFL6			(*) 9
	IL6			(*) 10
	TFF1			(*) 11

Candidates	BRCA1	(*)	12
	IGFBP1	(*)	13
	ATP5G3	(*)	14
	MUC5B	(*)	15
	SEMA6A	(*)	16
	FUT3	(*)	17
	MET	(*)	18
	GPR64	(*)	19
	ZBED1	(*)	20
	CPNE8	(*)	21
	SCGB1D2	(*)	22
	PAX8	(*)	23
	KLK7	(*)	24
	STC1	(*)	25
	CAPN2	(*)	26
	TPX2	(*)	27
	GAD1	(*)	28
	DKK4	(*)	29
	GCNT3	(*)	30
	CYP4B1	(*)	31
	LGALS4	(*)	32
	C1orf173	(*)	33
	C10orf116	(*)	34
	PBX1	(*)	35
	KGFLP2	(*)	36
	SENP8	(*)	37
	BCL2	(*)	38
	ADCYAP1R1	(*)	39
	TSPAN8	(*)	40
	LIN28B	(*)	41
	SERPINA5	(*)	42
	IGJ	(*)	43
	IGKC	(*)	44

4.3.2 SMOTE-Random Forest




Figure 4.11: Gene Optimization for SMOTE-Random Forest Classifier

In the SMOTE-Random Forest classifier, the F1-score stabilizes at around 0.7 when we reach 18 genes added, hence the optimal number of genes used will be $n=28+18=46$. The added genes are: TFF1, HNF1B, TFF3, LGALS4, SLC3A1, WT1, KLK7, TPX2, CYP2C18, GAD1, IGFBP1, CAPN2, FUT3, DKK4, C1orf173, GCNT3, C10orf116 and MUC5B.

Table 4.14: Gene Profile of Optimal Set in SMOTE-Random Forest Workflow

Set	Genes	PrOTYPE	SPOT	Optimal Set	Candidate Rank
	COL11A1	v		(*)	
	CD74	v		(*)	
	CD2	v		(*)	
	TIMP3	v		(*)	
	LUM	v		(*)	
	CYTIP	v		(*)	
	COL3A1	v		(*)	

Base	THBS2	v	(*)	
	TCF7L1	v	v	(*)
	HMGA2	v		(*)
	FN1	v		(*)
	POSTN	v		(*)
	COL1A2	v		(*)
	COL5A2	v		(*)
	PDZK1IP1	v		(*)
	FBN1	v		(*)
	HIF1A		v	(*)
	CXCL10		v	(*)
	DUSP4		v	(*)
	SOX17		v	(*)
	MITF		v	(*)
	CDKN3		v	(*)
	BRCA2		v	(*)
	CEACAM5		v	(*)
	ANXA4		v	(*)
	SERPINE1		v	(*)
	CRABP2		v	(*)
	DNAJC9		v	(*)
	TFF1			(*) 1
	HNF1B			(*) 2
	TFF3			(*) 3
	LGALS4			(*) 4
	SLC3A1			(*) 5
	WT1			(*) 6
	KLK7			(*) 7
	TPX2			(*) 8
	CYP2C18			(*) 9
	GAD1			(*) 10
	IGFBP1			(*) 11
	CAPN2			(*) 12
	FUT3			(*) 13

Candidates	DKK4	(*)	14
	C1orf173	(*)	15
	GCNT3	(*)	16
	C10orf116	(*)	17
	MUC5B	(*)	18
	ATP5G3	(*)	19
	PAX8	(*)	20
	IL6	(*)	21
	GPR64	(*)	22
	CPNE8	(*)	23
	PBX1	(*)	24
	STC1	(*)	25
	MET	(*)	26
	IGKC	(*)	27
	EPAS1	(*)	28
	TSPAN8	(*)	29
	SEMA6A	(*)	30
	EGFL6	(*)	31
	TP53	(*)	32
	CYP4B1	(*)	33
	KGFLP2	(*)	34
	BRCA1	(*)	35
	LIN28B	(*)	36
	SERPINA5	(*)	37
	BCL2	(*)	38
	SCGB1D2	(*)	39
	ZBED1	(*)	40
	SENP8	(*)	41
	ADCYAP1R1	(*)	42
	MAP1LC3A	(*)	43
	IGJ	(*)	44

4.3.3 Two-Step




Figure 4.12: Gene Optimization for Two-Step Classifier

Table 4.15: Gene Profile of Optimal Set in Two-Step Workflow

Set	Genes	PrOTYPE	SPOT	Optimal Set	Candidate Rank
	COL11A1	v		(*)	
	CD74	v		(*)	
	CD2	v		(*)	
	TIMP3	v		(*)	
	LUM	v		(*)	
	CYTIP	v		(*)	
	COL3A1	v		(*)	
	THBS2	v		(*)	
	TCF7L1	v	v	(*)	
	HMGA2	v		(*)	

Base	FN1	v	(*)	
	POSTN	v	(*)	
	COL1A2	v	(*)	
	COL5A2	v	(*)	
	PDZK1IP1	v	(*)	
	FBN1	v	(*)	
	HIF1A		v	(*)
	CXCL10		v	(*)
	DUSP4		v	(*)
	SOX17		v	(*)
	MITF		v	(*)
	CDKN3		v	(*)
	BRCA2		v	(*)
	CEACAM5		v	(*)
	ANXA4		v	(*)
	SERPINE1		v	(*)
	CRABP2		v	(*)
	DNAJC9		v	(*)
	CYP2C18		(*)	1
	MUC5B		(*)	2
	HNF1B		(*)	3
	IL6		(*)	4
	SLC3A1		(*)	5
	EGFL6		(*)	6
	WT1		(*)	7
	ZBED1		(*)	8
	MET		(*)	9
	SENP8		(*)	10
	KLK7		(*)	11
	TFF3		(*)	12
	CPNE8		(*)	13
	STC1		(*)	14
	GAD1		(*)	15
	LIN28B		(*)	16

Candidates	IGJ	(*)	17
	DKK4	(*)	18
	EPAS1	(*)	19
	GCNT3	(*)	20
	SCGB1D2	(*)	21
	CYP4B1	(*)	22
	C1orf173	(*)	23
	IGFBP1	(*)	24
	TPX2	(*)	25
	SEMA6A	(*)	26
	ATP5G3	(*)	27
	SERPINA5	(*)	28
	FUT3	(*)	29
	C10orf116	(*)	30
	KGFLP2	(*)	31
	ADCYAP1R1	(*)	32
	TP53	(*)	33
	PBX1	(*)	34
	GPR64	(*)	35
	LGALS4	(*)	36
	CAPN2	(*)	37
	BCL2	(*)	38
	MAP1LC3A	(*)	39
	TSPAN8	(*)	40
	TFF1	(*)	41
	PAX8	(*)	42
	BRCA1	(*)	43
	IGKC	(*)	44

4.4 Test Set Performance

Now we'd like to see how our best methods perform in the confirmation and validation sets. The class-specific F1-scores will be used.

The top 2 methods are the sequential and SMOTE-Random Forest classifiers. We can test 2 additional methods by using either the full set of genes or the optimal set of genes for both of these

classifiers.

4.4.1 Confirmation Set

Table 4.16: Evaluation Metrics on Confirmation Set Models

Method	Metric	Overall	Histotypes				
			HGSC	CCOC	ENOC	LGSC	MUC
Sequential, Full Set	Accuracy	0.834	0.869	0.960	0.894	0.977	0.967
	Sensitivity	0.655	0.948	0.853	0.462	0.308	0.704
	Specificity	0.928	0.719	0.974	0.980	0.990	0.979
	F1-Score	0.664	0.905	0.831	0.590	0.348	0.644
	Balanced Accuracy	0.792	0.834	0.913	0.721	0.649	0.841
	Kappa	0.665	0.697	0.808	0.535	0.336	0.627
Sequential, Optimal Set	Accuracy	0.821	0.865	0.949	0.879	0.983	0.967
	Sensitivity	0.633	0.953	0.840	0.396	0.385	0.593
	Specificity	0.923	0.697	0.963	0.974	0.995	0.984
	F1-Score	0.659	0.902	0.792	0.519	0.476	0.604
	Balanced Accuracy	0.778	0.825	0.902	0.685	0.690	0.788
	Kappa	0.635	0.684	0.763	0.457	0.468	0.587
SMOTE-Random Forest, Full Set	Accuracy	0.843	0.871	0.969	0.896	0.980	0.970
	Sensitivity	0.659	0.962	0.867	0.453	0.308	0.704
	Specificity	0.928	0.697	0.982	0.983	0.994	0.982
	F1-Score	0.682	0.907	0.867	0.589	0.381	0.667
	Balanced Accuracy	0.793	0.829	0.925	0.718	0.651	0.843
	Kappa	0.677	0.697	0.849	0.535	0.371	0.651
SMOTE-Random Forest, Optimal Set	Accuracy	0.851	0.876	0.966	0.904	0.981	0.975
	Sensitivity	0.695	0.957	0.853	0.500	0.385	0.778
	Specificity	0.932	0.719	0.981	0.983	0.994	0.984
	F1-Score	0.715	0.910	0.853	0.631	0.455	0.724
	Balanced Accuracy	0.813	0.838	0.917	0.742	0.689	0.881
	Kappa	0.697	0.710	0.834	0.580	0.445	0.711
Two-Step, Full Set	Accuracy	0.843	0.869	0.960	0.904	0.975	0.978
	Sensitivity	0.672	0.948	0.853	0.509	0.308	0.741
	Specificity	0.930	0.719	0.974	0.981	0.989	0.989
	F1-Score	0.689	0.905	0.831	0.635	0.333	0.741
	Balanced Accuracy	0.801	0.834	0.913	0.745	0.648	0.865
	Kappa	0.683	0.697	0.808	0.584	0.321	0.729
Two-Step, Optimal Set	Accuracy	0.840	0.869	0.958	0.899	0.977	0.977
	Sensitivity	0.645	0.955	0.853	0.481	0.231	0.704
	Specificity	0.928	0.706	0.972	0.981	0.992	0.989
	F1-Score	0.669	0.906	0.826	0.611	0.286	0.717
	Balanced Accuracy	0.786	0.830	0.913	0.731	0.611	0.846
	Kappa	0.673	0.695	0.802	0.557	0.275	0.705




Figure 4.13: Gene Optimized Workflows Per-Class Metrics in Confirmation Set

Confusion Matrices for Confirmation Set Models




Figure 4.14: Confusion Matrices for Confirmation Set Models

4.4.1.1 Sequential, Full

ROC Curves for Sequential, Full Set Model in Confirmation Set




Figure 4.15: ROC Curves for Sequential Full Model in Confirmation Set

4.4.1.2 Sequential, Optimal

ROC Curves for Sequential, Optimal Set Model in Confirmation Set

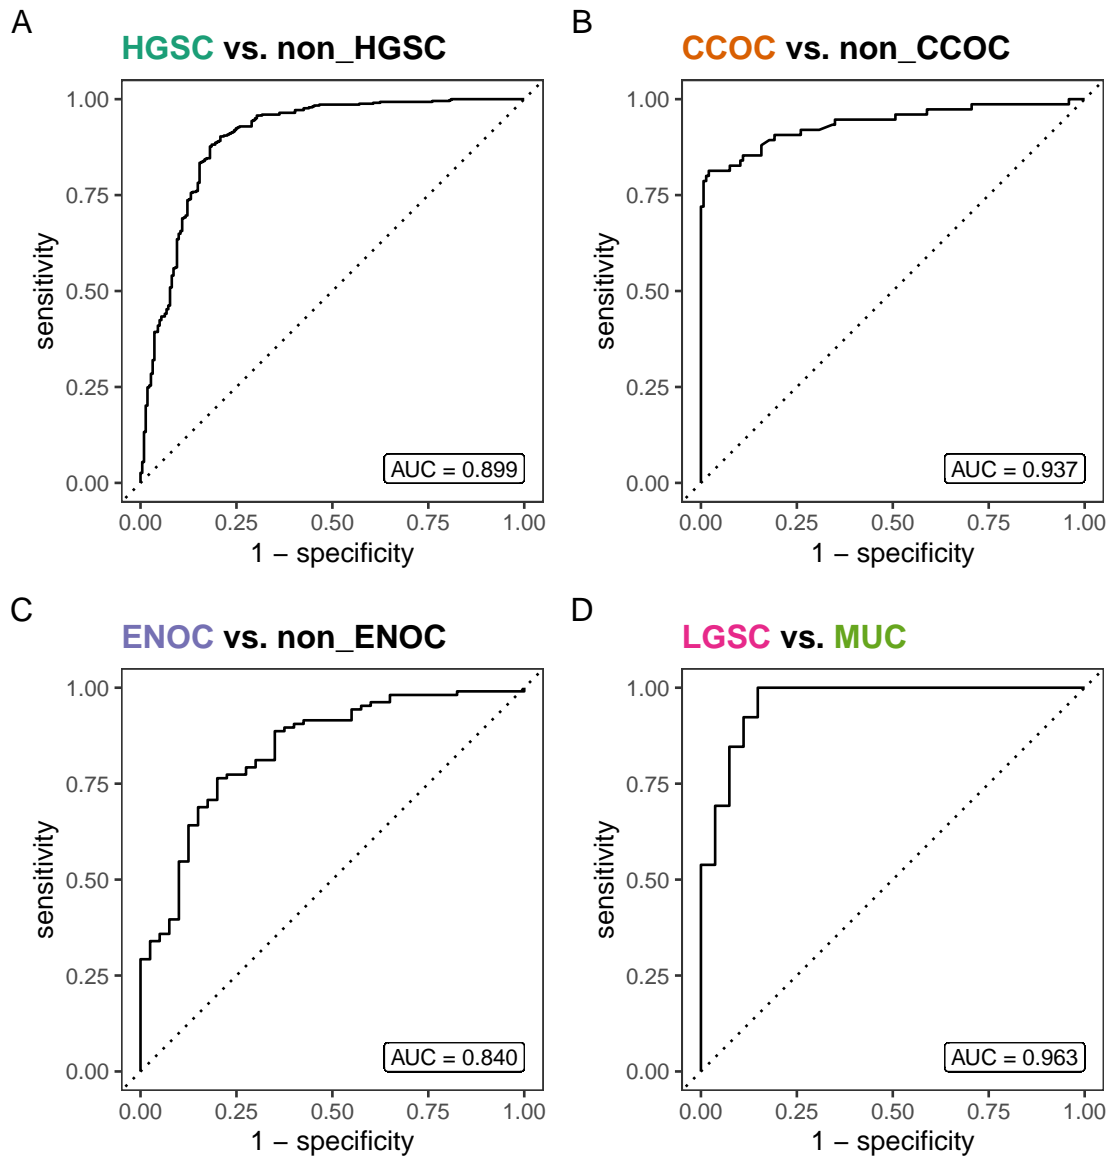


Figure 4.16: ROC Curves for Sequential, Optimal Model in Confirmation Set

4.4.1.3 SMOTE-Random Forest, Full




Figure 4.17: ROC Curves for SMOTE-Random Forest, Full Set Model in Confirmation Set

4.4.1.4 SMOTE-Random Forest, Optimal




Figure 4.18: ROC Curves for SMOTE-Random Forest, Optimal Set Model in Confirmation Set

4.4.1.5 Two-Step, Full

ROC Curves for Two-Step, Full Set Model in Confirmation Set

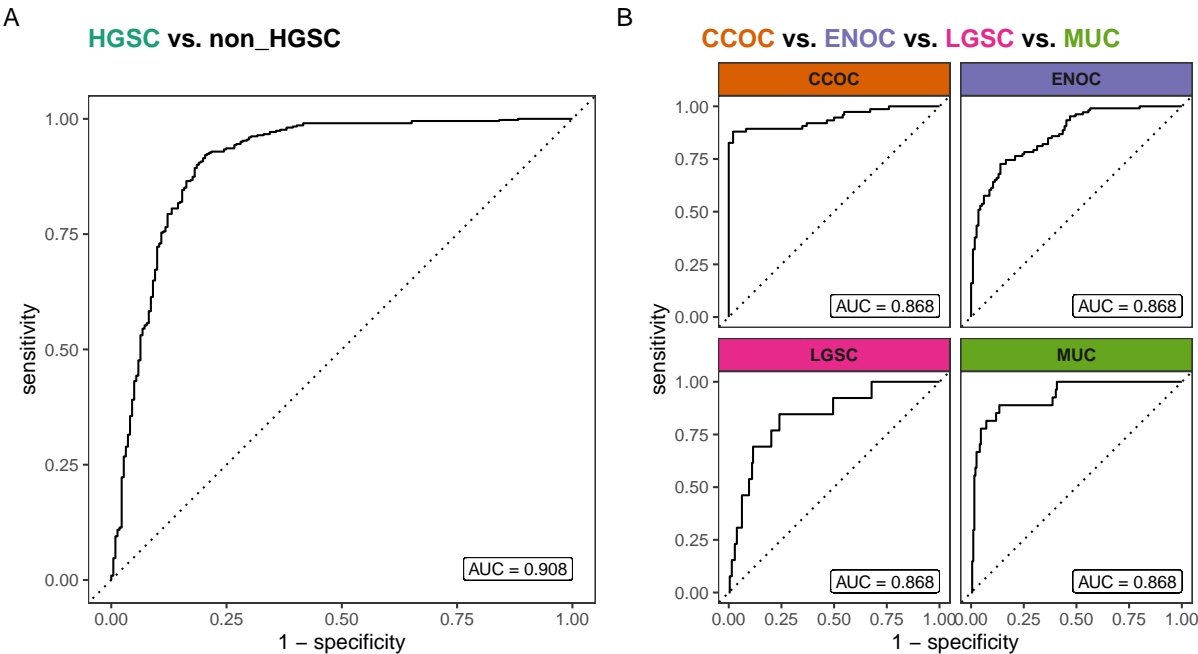


Figure 4.19: ROC Curves for Two-Step Full Model in Confirmation Set

4.4.1.6 Two-Step, Optimal

ROC Curves for Two-Step, Optimal Set Model in Confirmation Set

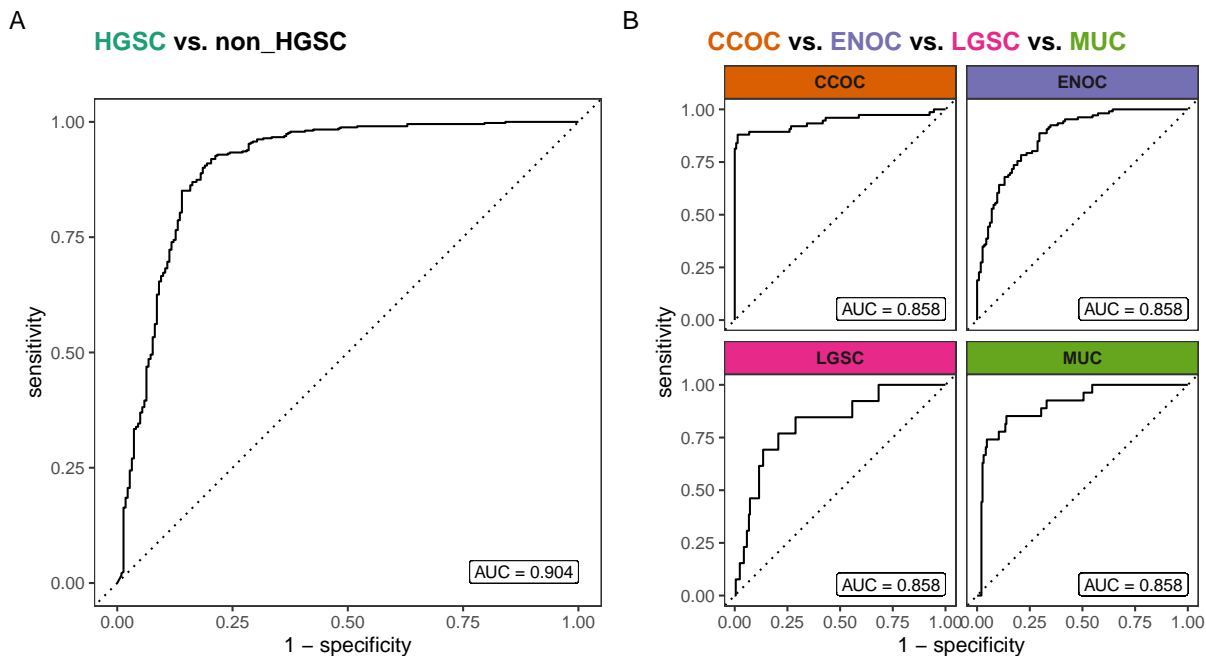


Figure 4.20: ROC Curves for Two-Step Optimal Model in Confirmation Set

4.4.2 Validation Set

Table 4.17: Evaluation Metrics on Validation Set Model, SMOTE-Random Forest, Optimal Set

Metric	Overall	Histotypes				
		HGSC	CCOC	ENOC	LGSC	MUC
Accuracy	0.890	0.907	0.971	0.952	0.972	0.979
Sensitivity	0.774	0.926	0.937	0.714	0.444	0.846
Specificity	0.955	0.851	0.974	0.984	0.983	0.983
F1-Score	0.731	0.937	0.851	0.777	0.390	0.698
Balanced Accuracy	0.864	0.889	0.955	0.849	0.714	0.914
Kappa	0.748	0.761	0.835	0.750	0.376	0.688




Figure 4.21: SMOTE-Random Forest Per-Class Metrics in Validation Set

Confusion Matrix for Validation Set Model

SMOTE–Random Forest, Optimal Set

Prediction	HGSC	617	5	21	8	0
	CCOC	13	74	6	1	1
	ENOC	10	0	75	0	3
	LGSC	15	0	0	8	0
	MUC	11	0	3	1	22
		HGSC	CCOC	ENOC	LGSC	MUC
		Truth				

Figure 4.22: Confusion Matrix for Validation Set Model

ROC Curve for SMOTE–Random Forest, Optimal Set Model in Validation Set

AUC = 0.959




Figure 4.23: ROC Curves for SMOTE-Random Forest, Optimal Set Model in Validation Set

Subtype Prediction Summary among Predicted HGSC Samples

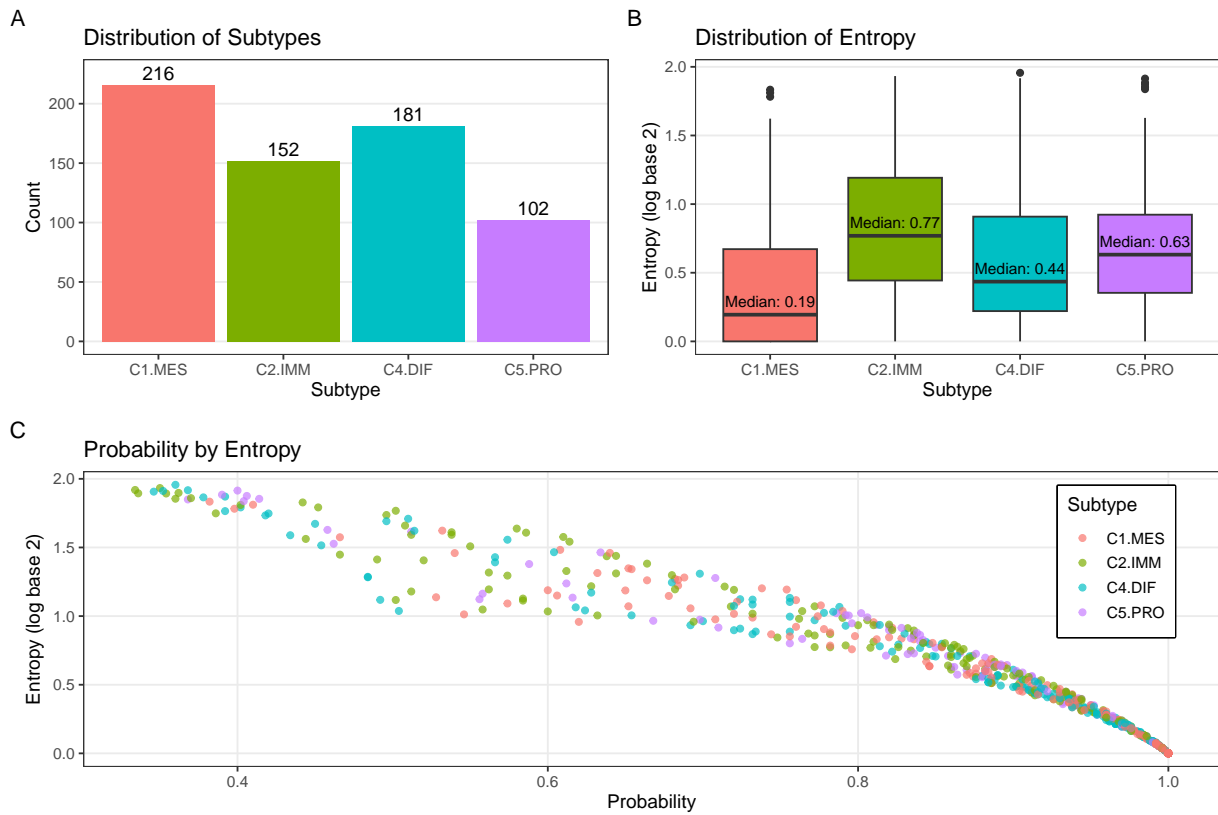


Figure 4.24: Subtype Prediction Summary among Predicted HGSC Samples

References

Talhouk, Aline, Stefan Kommoss, Robertson Mackenzie, Martin Cheung, Samuel Leung, Derek S. Chiu, Steve E. Kalloger, et al. 2016. “Single-Patient Molecular Testing with NanoString nCounter Data Using a Reference-Based Strategy for Batch Effect Correction.” Edited by Benjamin Haibe-Kains. *PLOS ONE* 11 (4): e0153844. <https://doi.org/10.1371/journal.pone.0153844>.